

Peer Networks and Educational Migration: Experimental Evidence from India

Merve Demirel*

September 16, 2026

Preliminary and incomplete — please do not cite or circulate.

[Click here for the latest version.](#)

Abstract

Migration in developing countries is low despite large differences in economic opportunity across space. This paper asks whether the non-pecuniary costs of migration hold prospective college students back from migrating for college, and whether access to a migrant peer can reduce them. I run a randomized controlled trial with 1,682 final-year secondary students across 52 schools in Maharashtra, India, offering half a phone-based mentorship program that pairs each student with a university student who migrated for college. At baseline, the average student would forgo a 129% increase in expected post-graduation earnings rather than study away from home, a penalty that rises steeply with distance, and beliefs about the difficulty of moving, safety, parental approval, and happiness away from home account for about 61.3% of it. Mentorship shifts choices concentrated at margin: applications and moves to the closest destination rise, the pre-specified educational migration index rises by 0.16 standard deviations, and the effect is concentrated among men. Non-pecuniary beliefs about the distant cities do not improve.

*PhD student, Institute for International Economic Studies (IIES), Stockholm University. Email: merve.demirel@iies.su.se. The trial is registered in the AEA RCT Registry (AEARCTR-0016735), and pre-analysis plans for the midline and endline surveys were filed before the respective data were collected. The study received ethical approval from the DAI Research & Advisory Services Institutional Review Board (approval no. PN-MH-01/2025-2026) and the Swedish Ethical Review Authority (Etikprövningsmyndigheten, approval no. 2025-04262-01). I gratefully acknowledge financial support from the Royal Swedish Academy of Sciences (Kungliga Vetenskapsakademien) and the Weiss Fund for Research in Development Economics.

1 Introduction

Income differences across space are among the largest and most persistent in developing countries, yet few people move (Gollin et al., 2014; Lagakos, 2020; Kone et al., 2018). A large literature attributes this to a large disutility of migration itself: estimates of the compensation required to leave home are large, whether inferred from observed choices or elicited directly (Kennan and Walker, 2011; Bryan and Morten, 2019; Lagakos et al., 2023; Koşar et al., 2022). Less is known about what this disutility consists of, or whether it can be changed.

For many people, the first migration decision comes with the choice of where to study. Higher-education opportunities are sharply uneven across space: a student who moves gains access to better programs and starts a career in a higher-wage labor market. Whether they stem from high costs of migration, or misperceptions about the returns to moving itself, barriers to migration are therefore barriers to college access, and their cost accrues throughout a working life. How large are these barriers, what do they consist of, and can access to a peer who has already made a similar move reduce them, thereby changing whether and where students choose to study?

I study these questions in the context of college choice in Maharashtra, India, where students decide whether and where to continue their education, and in doing so largely determine where they build their adult lives. The median student in my sample expects studying at a college away from home to lead to higher earnings net of living costs, by about \$2,300 in net present value over a working life. Among the 56% of students who expect any positive return, the median expected gain is roughly \$15,000. Yet, a majority name their home as their most preferred study destination (Section 2). Whether this represents student's preferences, misperceptions, constraints or real costs of migration determines how the spatial income gaps should be understood: if the unwillingness to move reflects informed preferences, the gap consists of compensating differentials and staying home is efficient. Otherwise, the same gap can be characterized as unrealized welfare, and the question becomes whether there is room for policy interventions to help students realize such gains.

These explanations cannot be separated on choice data alone: the same decision to stay home is consistent with a student who knows what life away from home would be like and dislikes it, and with one whose beliefs about it are malleable (Manski, 2004). I therefore proceed in two steps. First, I measure the unwillingness to move directly: in a discrete choice experiment that holds feasibility fixed and an incentivized college-rating exercise, I estimate how much of their expected earnings students will forgo to study

closer to home. I then decompose the observed wedge between the returns students expect and the choices they make into non-pecuniary expectations: I elicit what students expect if studying in each destination (safety, social ties, emotional wellbeing) and ask how much of the wedge these beliefs account for. At a second step, to test whether the wedge reflects barriers that can be addressed by new information, I run a randomized controlled trial that expands students' access to the people best placed to offer such information: I randomly offer half of 1,682 secondary school students across 52 schools in Maharashtra a place in *Campus Trail*, a phone-based program that matches each student with a current university student who left their hometown to study in Pune or Mumbai. The mentor-mentee pairs speak by phone over the months before college application deadlines. I estimate effects on realized application, enrollment, and migration decisions at endline, and on destination-specific beliefs.

In the discrete choice experiment, the average student would forgo a 129% increase in their own expected post-graduation earnings rather than study away from home; the penalty rises with distance, from 39% for a move to the district capital to more than double their earnings for the large cities. The design of the choice experiment rules out the constraint story as the only explanation: even if educational migration is feasible, students attach a large disutility to migration. Because the design also holds pecuniary returns fixed, it isolates the non-pecuniary aspects of this disutility. Eliciting beliefs about migrating and studying in these cities, I find that a small set, about safety, proximity to family, and happiness, accounts for a large share. The pecuniary returns students expect from moving are not large enough to compensate, and most plan to study at home. The pattern is gendered, with women expecting larger declines in safety, approval, and happiness studying away from home, and therefore require a larger salary increase to leave. I then examine the information environment these beliefs are built in, and find that students have few sources of information to verify these beliefs against: at baseline, 29% of students report having no source of college advice at all, and 40% have no contacts in the destination cities they were asked about.

Mentorship shifts students' applications and realized moves. At endline, a pre-specified migration index for any move away from home rises by 0.16 standard deviations. Treated students are three percentage points more likely to have applied for further study and five points more likely to be enrolled or planning to enroll; applications to colleges away from home rise by four percentage points and moves away from home by three, both imprecisely estimated.

The moves the program induces are the shortest ones, where the perceived non-pecuniary costs of migration is smallest. Consistent with the stated-preference evidence,

the additional applications and moves go to the district capital, by five and two percentage points, while applications and migration to the distant cities, where the non-pecuniary penalty is largest, do not increase. The effect is concentrated among students living outside a district capital at baseline, three quarters of the sample, whose index rises by 0.21 standard deviations, and among men: male students' applications away from home rise by ten percentage points while women's do not move. These are precisely the students and the moves for which the non-pecuniary cost of moving is smallest.

I then investigate how mentorship changed beliefs about the destinations. At endline, treated students expect the district capital to be cheaper relative to home, and among men. Importantly, non-pecuniary beliefs about the distant cities are unchanged for men and lower for women; information, search effort, and trait measures (gender attitudes, aspirations, self-efficacy) do not change. I read this as consistent with the large non-pecuniary costs students perceive at baseline being broadly accurate rather than misperceived. If these beliefs were largely driven by misperceptions, contact with someone who made the move could correct the beliefs, even if the resulting choices do not move. I find evidence for the contrary, which is consistent with real non-pecuniary costs of migration, large enough to not be affected by a light-touch peer mentorship intervention, as a key barrier for educational mobility.

The paper's first contribution is to the literature on why people migrate so little despite large spatial income gaps. This literature consistently estimates moving costs far larger than the observable monetary costs of a move, with consequences for productivity, the allocation of labor, and human capital (Sjaastad, 1962; Kennan and Walker, 2011; Bryan and Morten, 2019; Lagakos et al., 2023; Ishimaru, 2025; Fu et al., 2022; Fabre, 2024). What these costs consist of is largely unknown, because they are difficult to observe, and the few studies that elicit them directly do so from hypothetical choices (Koşar et al., 2022; Baláž et al., 2016; Batista et al., 2025) or infer them from social ties (Sahai and Bailey, 2022; Koenen and Johnston, 2024). Reviews of this literature identify non-monetary costs and attachment to home as the least understood part of the decision (McKenzie, 2022). Measured directly, I find that this disutility is large, that a handful of beliefs about life away from home account for most of it, and that it is hard to move: contact with a migrant peer changes choices only where the penalty is already small, and does not improve beliefs about the distant destinations. The experiment also adds to the evidence on networks and migration. People with connections at the destination migrate more (Munshi and Rosenzweig, 2016; Blumenstock et al., 2025; Aydemir et al., 2022; Koenen and Johnston, 2024; Green, 2024), but connections are chosen, so these estimates

mix the effect of a connection with selection into having one. The mentorship program connects students at random, and in doing so joins a small set of experiments that use exposure to lived migration experience to learn about migration behavior: Barnett-Howell et al. (2025) pair rural job-seekers in Kenya with urban migrant mentors, Beber et al. (2026) run a mentoring program for prospective migrants in Senegal, and Mespl  -Soms and Nilsson (2023) expose young people in rural Mali to migrant role models, with mentoring raising migration while role models move aspirations more than the desire to migrate. I test a similar intervention in an educational migration setting and show that such exposure affects beliefs about the experiential costs of moving, not only the monetary returns.

The paper also contributes to the literature on the geography of educational opportunity. Access to education depends on where students live: building schools raises attainment (Duflo, 2001), opening colleges raises enrollment, especially among women (Caner et al., 2024; Sekhri et al., 2024; Elsayed and Shirshikova, 2023), and demand for existing institutions falls steeply with distance, enough for spatial frictions to produce regional and demographic gaps in attainment (Card, 1995; Acton et al., 2025; Alba-Vivar, 2024). Much of this literature evaluates supply-side policies that bring education closer to students, by changing where institutions are or what it costs to commute to them. A newly opened college within commuting distance may increase enrollment, but the education it offers is unlikely to match that of institutions established in the cities, and it leaves the student in the same labor market if the constraints to migration bind, since graduates largely enter the labor market where they studied (Conzelmann et al., 2025). If so, there is a case for policies that move students to education rather than education to students. Consistent with this, Ghose (2026) estimates that reducing the cost of moving for education in India raises welfare by more than reducing the cost of moving for work, because the student who moves for college need not move again for a job. My experiment evaluates a demand-side intervention that works exactly on this margin. In my setting, students can continue studying near home, but the options there are few and less selective. Instead, considering options a few hours of travel distance away, a student can reach more than tenfold seats, with the largest gains among the programs applicants compete for most (Section 2). Reaching those seats means living away from home, a majority of students are not willing to consider these options at baseline.

The paper also speaks to work on gendered mobility. Distance, travel costs, and safety shape women’s schooling and labor-market choices differently from men’s (Muralidharan and Prakash, 2017; Borker, 2021; Le Barbanchon et al., 2021; Farre and Ortega, 2021; Field and Vyborny, 2022; Kapoor and Gade, 2024; Alba-Vivar, 2024; Cheema et al.,

2025; Han et al., 2025): a bicycle or a safer route is enough to change where girls enroll, and the same frictions reappear at each step between school and work. I show that they extend from the daily commute to the high-stakes, longer-term decision of where to study. Women would require 227% increase in expected earnings to study away, against 72% for men. They also perceive larger decreases in safety, parental approval and other non-pecuniary outcomes if they were to study away, and contact with a migrant peer does not improve them. If these frictions are subjective but accurate, information alone will not close the gender gap in mobility: I consider my findings as evidence that these perceived decreases are broadly accurate, rather than driven by misperceptions.

Finally, the paper adds to the measurement of subjective expectations in education. Beliefs about earnings and academic success shape human-capital choices (Jensen, 2010; Wiswall and Zafar, 2015; Giustinelli, 2023). A smaller body of work extends the elicitation to non-pecuniary returns: Boneva and Rauh (2017) show that beliefs about job satisfaction, social life, and parental approval explain much of the socio-economic gap in university enrollment intentions, Delavande and Zafar (2024) find that non-pecuniary factors such as a school's religious ideology, rather than expected earnings, drive university choice in Pakistan, and Clerici et al. (2024) find that family approval matters for the occupational choices of young adults in Uganda. I focus on the beliefs specific to the migration decision: what daily life in a different city would look like. I elicit these for each destination, estimate the weights students attach to them, and experimentally test whether they move. The intervention itself builds on evidence that mentors and role models shift schooling and early career choices (Bettinger et al., 2012; Carrell and Sacerdote, 2017; Castleman and Page, 2015; Beaman et al., 2012; Riley, 2024; Breda et al., 2023; Alfonsi et al., 2022; Hoxby and Turner, 2013). A consistent lesson from this work is that students facing complex decisions need personalized and intensive support, including information, encouragement, and practical help with various application procedures. The intervention applies that lesson to a distinct margin of college access: barriers generated and amplified by spatial distance.

The remainder of the paper proceeds as follows. Section 2 describes the setting and sample. Section 3 outlines the conceptual framework. Section 4 documents the size of the migration penalty, the beliefs behind it, and the information environment in which those beliefs form. Section 5 describes the experiment and data, and Section 6 describes the empirical strategy. Section 7 presents the treatment effects of the intervention. Section 8 concludes.

2 Setting and sample

2.1 Higher education and migration in Maharashtra

The study follows 12th-grade students sampled across 52 schools in two districts in Maharashtra, as they decide where to continue their education. Throughout the paper, I use *home* for the student's own *taluka*, or sub-district,¹ and *away* for any destination outside it. The mentors are college students in Pune and Mumbai² regions, two of Maharashtra's main higher-education and labor-market destinations (Figure 1).

Higher-education supply varies sharply across Maharashtra (Appendix Figure A1), with Pune and the Mumbai region, where the mentors study, constituting the main higher-education hubs of Maharashtra. According to All India Survey on Higher Education (AISHE) data from 2022, these two regions together host a quarter of Maharashtra's universities and affiliated colleges, and a third of the students enrolled in them. They also absorb a disproportionate share of the state's educational migrants: together, they account for 28 % of within-Maharashtra educational migrants, the largest concentration of any destination cluster in the state.

For a prospective college student in the study sample, colleges a few hours of travel away therefore offer drastically wider access to higher education. Figure 2 makes this concrete for one origin: it traces the number of engineering and business-program seats a general-category applicant from Sangamner taluka, in Ahilyanagar, could access through the centralized admissions system at each point of the entrance-exam ability distribution. Even the strongest such applicant can reach only about 250 engineering and 100 business-program seats within their *home*, and fewer than 1,000 engineering and 400 business seats anywhere in Ahilyanagar; the out-of-district set is more than an order of magnitude larger: roughly 22,000 and 5,500 seats, respectively.

Seat counts understate the out-of-district advantage in two ways. First, high-demand programs are concentrated in Pune and the Mumbai region: the distribution of admission cutoffs there lies above that of the study districts (Figure A3), so the programs applicants compete for most are disproportionately in Pune and Mumbai. Second, the larger market is not offset by more applicants competing for it: accounting for the number of eligible candidates, out-of-district options still offer more seats per candidate, and the advantage grows with a student's achievement (Figure A2). For students in Ahilyanagar

¹A *taluka* is a sub-district administrative area comprising multiple villages and towns. I define a student's home as the taluka in which their school is located and treat it as the student's local geographic, labor-market, educational, and social environment. I reserve *out-of-district* for away destinations beyond the home district, such as Pune, Mumbai, and Nagpur, and use it only where the district boundary itself matters. The baseline survey modules measured home at the district level, so every baseline away option lies outside the home district (Appendix C).

²The Mumbai region comprises Mumbai, Mumbai Suburban, and Thane districts.

and Sangli, access to this part of the higher-education market therefore requires studying away from home.

The economic gradient reinforces the educational one. Pune and the Mumbai region offer access not only to more selective colleges, but also to larger and higher-paying labor markets. This matters beyond the very top of the academic distribution: even holding college quality fixed, studying in these destinations may help students become locally connected to better jobs through access to internships, job-search networks, or local employers. Per-capita incomes in Pune and the Mumbai region are 60–90% higher than in the study districts,³ a gap that remains even after adjusting for cost of living differences. Household survey data point in the same direction: in the India Human Development Survey, which is not representative at the district level, per-capita incomes in these two regions are 105% higher and consumption 57% higher than in the rest of the state after adjusting for spatial price differences, and among households with a college graduate the income gap is still 51% (Appendix Table A1). These observational differences reflect both selection and place effects, so they should not be interpreted as the causal return a marginal student would obtain by migrating; they do, however, indicate the magnitude of the gap a prospective educational migrant is weighing.

Despite this gradient, migration for education in Maharashtra is uncommon and falls steeply with distance: according to the 2011 Census, only 1.35 % of within-state migrants moved for education, and 57.5 % of those educational migrants stayed within their own district.⁴ A fragmented admissions system comprising multiple entrance exams that vary by program and region, and state- and education-region-level quotas that reserve seats for local applicants, reinforce the incentive to stay close to home. In this study's sample, 9.0% of graduates in the control group migrate within district for education, 11.8% across district, and virtually none migrate across state borders.

2.2 Sample

The sample consists of 1,682 final-year secondary-school students across 52 schools. Sixty percent of students are female, 68% study a science stream, 36% belong to the general (non-reserved) caste category, and the median student is seventeen. About 70% of students have fathers who have completed at least lower-secondary education. Relative to

³Source: Directorate of Economics and Statistics, Planning Department, Government of Maharashtra, *District Domestic Product of Maharashtra 2011–12 to 2023–24* (<https://mahades.maharashtra.gov.in>).

⁴The patterns are similar at the national level. While 38 % of the Indian population moved within the country, only 1.1 % of all migrants moved for education, compared to 50 % for marriage and 10.4 % for employment or business; and although 20.8 % of young Indians are in higher education (Ministry of Human Resource Development, 2013), only 3.3 % migrate for education. Among those who do move, 53.5 % stay within their own district and only 13.9 % cross state lines. Migration is also strongly gendered; Section 4 returns to gender heterogeneity in the context of the migration penalty.

district peers, the sample is mildly disadvantaged on school environment: study-sample schools are more likely to be rural, and have fewer teachers and classrooms per student (Appendix Table A16).

Students in the sample are preparing for the Higher Secondary Certificate (HSC) board exam at the end of twelfth grade, which is required for graduation and serves as one admission route: many institutions admit on the basis of HSC marks alone. Most students additionally sit for one or more centralized entrance exams. In Maharashtra, the MHT-CET admits to state engineering, pharmacy, and management programs; the national JEE admits to the elite IITs and NITs; NEET admits to medical programs; and a range of other exams admit to law, agriculture, and design programs. Vocational and diploma tracks have their own admission systems. Coaching, as after-school classes during higher secondary school or as full-time after, is a standard part of preparation across all tracks: 60% of students report already attending or planning to attend coaching classes at baseline.

Two features of the sample matter for interpreting the experiment. First, educational migration is academically feasible for the average student in the sample. The typical study-sample student lies around the 25th percentile of students placing into Pune and the Mumbai region on the MHT-CET, and Appendix Figure A9 shows that HSC results in study-sample schools are lagging behind district and state averages only slightly. Second, migration is financially feasible on average. Students report being able to budget around \$76 per month for college at baseline, which is comparable to the budgets of a representative sample of students in Pune (Mangal, 2024), and at the lower end of the mentor applicant pool (Appendix B).

The sample therefore sits on the relevant margin for the experiment. Students expect substantial pecuniary gains from studying in the distant cities elicited in the survey compared to home: they expect higher post-graduation returns (16%) on average; 56 % of the sample report positive returns net of living costs for at least one of the two elicited cities, so this is a lower bound on the share perceiving a paying destination somewhere. Among them, the median stated differential is roughly \$15,100 (INR 14 lakh) in present-value terms, with a long right tail (mean \$30,700; INR 29 lakh); over the full sample, the median differential is about \$2,300 (INR 2.1 lakh).⁵ Yet when asked directly for their most preferred option, 54 % of these same students name their own district, and only

⁵For each elicited city, the pecuniary differential combines the stated monthly salary (bracket midpoints) and monthly living-cost gaps relative to home from the baseline city module. The present value assumes the living-cost differential is borne over four years of study and the salary differential earned over the following forty years, discounted at 5 % annually; each student's best destination is used. The positive-return share is 55 % when the salary differential is weighted by the student's stated probability of working full-time at age 25. In a second conservative check, expected salary after studying away is $p_{ic}^{\text{stay}} w_{ic} + (1 - p_{ic}^{\text{stay}}) w_{ih}$, where p_{ic}^{stay} is the student's destination-specific probability of remaining there after graduation; the positive-return share is 50 %. Applying both probability adjustments yields 49 %. USD at 93.75 INR.

26 % name a destination they themselves expect to pay off.⁶ The rest of the paper asks whether this gap reflects non-pecuniary migration costs, and whether access to migrant peers can change the beliefs and choices that sustain them.

3 Conceptual framework

I conceptualize students' college migration decisions as a destination-choice problem. Student i considers a set of destinations c and chooses the option that maximizes utility:

$$U_{ic} = f(V_{ic}, N_{ic}).$$

Here V_{ic} is the student's perceived net *pecuniary value* of attending college in destination c : the present value of expected post-graduation earnings net of expected pecuniary costs, including tuition fees, travel, housing, and living expenses. N_{ic} is the student's perceived net *non-pecuniary value* of attending college in c , bundling safety, parental approval, freedom, social ties, mental wellbeing, and other lived-experience factors. Utility is increasing in both arguments.

The objects V_{ic} and N_{ic} are subjective. They summarize the values that enter student i 's decision, combining beliefs about each destination with the weights students place on those expected conditions. Two students may therefore assign different V_{ic} or N_{ic} to the same destination: they may hold different beliefs about safety, costs, or returns, or they may place different weight on the same expected conditions.

Home is the default option and the natural reference point for the migration decision. A student migrates from home h to another destination c when

$$U_{ic} > U_{ih}.$$

The relevant comparison is therefore the destination-home utility gap:

$$\Delta U_{ic} \equiv U_{ic} - U_{ih},$$

with corresponding pecuniary and non-pecuniary gaps

$$\Delta V_{ic} \equiv V_{ic} - V_{ih}, \quad \Delta N_{ic} \equiv N_{ic} - N_{ih}.$$

I call ΔN_{ic} the *migration wedge*: the non-pecuniary value student i attaches to studying in

⁶Most preferred option from the baseline main survey's direct ranking question. A ratings-based alternative (the strictly highest-rated of the three elicited destinations on the 0–10 scale) gives a nearly identical 25 %.

c rather than at home. A positive wedge means the student expects to be better off away on non-pecuniary grounds; a negative wedge is a cost of migrating that the pecuniary gain ΔV_{ic} must overcome. Section 4 measures the wedge and finds it negative for the average student at every destination; I refer to its magnitude as the *migration penalty*.

A student may expect higher net pecuniary value from studying away from home, $\Delta V_{ic} > 0$, but still prefer home if the non-pecuniary gap is sufficiently negative. Conversely, a student may migrate because the pecuniary gain is large, because the non-pecuniary cost is small, or because the destination dominates home on both margins. Migration selection may therefore reflect high expected pecuniary returns away from home, low perceived migration costs, low sensitivity to those costs, or a combination of these forces.

The framework implies two conditions for the intervention to change observed migration choices. First, the intervention must raise the utility of at least one non-home destination relative to home:

$$\Delta U_{ic}^{post} > \Delta U_{ic}^{pre}.$$

Because utility is increasing in V and N , this can occur if the intervention raises the perceived or effective pecuniary value of the destination relative to home, ΔV_{ic} , or the perceived or effective non-pecuniary value of the destination relative to home, ΔN_{ic} .

Second, the shift must be large enough to reverse the student's preferred option. For a student who prefers home at baseline, this requires

$$\Delta U_{ic}^{pre} < 0 < \Delta U_{ic}^{post}.$$

Equivalently, the treatment-induced shift must exceed the student's baseline distance to the migration margin:

$$\Delta U_{ic}^{post} - \Delta U_{ic}^{pre} > -\Delta U_{ic}^{pre}.$$

The intervention may therefore change beliefs or reduce costs for many students, but observed migration changes only for students whose destination ranking is reversed.

Mechanisms The intervention can satisfy the first condition through several channels. First, peer conversations may change students' beliefs about destination-specific returns or costs. For example, students may revise their expectations about safety, loneliness, commuting, parental acceptance, living conditions, or the practical difficulty of navigating a destination.

Second, peer conversations may reduce effective migration costs directly. A peer may

help identify safe housing, explain how to move around the city, provide contacts, or make the destination easier to navigate. In this case, the intervention changes not only beliefs about migration costs but the costs themselves.

Third, the intervention may change the salience or perceived importance of particular costs, making the same expected difficulty feel more manageable. A change in ΔN_{ic} , for example, may therefore reflect belief updating about a fixed destination environment, a real reduction in migration costs through peer support, or a change in how strongly the student weights a given expected condition. Section 5 returns to how the intervention addresses each of these mechanisms.

Migration responses The ranking-reversal condition clarifies when changes in the migration calculus translate into observed migration. Students with ΔU_{ic}^{pre} close to zero require only a small shift in perceived or effective returns and costs for a non-home destination to become preferred. Students with large negative values of ΔU_{ic}^{pre} are unlikely to migrate unless the intervention produces a large shift. Students with large positive values of ΔU_{ic}^{pre} would migrate even without the intervention. Treatment effects on observed migration should therefore be concentrated among students close to the migration margin.

The framework motivates the empirical analysis that follows. Section 4 quantifies the non-pecuniary gap ΔN_{ic} and the beliefs behind it. Section 7 then estimates whether the intervention changes realized migration at endline, where the induced moves go, and which beliefs move, and for whom. Throughout, the ranking-reversal condition guides interpretation: belief movement need not translate into migration except for students close to the margin.

4 The non-pecuniary cost of migration

The setting section establishes the central puzzle: many students expect higher pecuniary value from destinations away from home, yet still prefer to study at home. In the conceptual framework in Section 3, this means that a negative non-pecuniary value differential, $\Delta N_{ic} < 0$, offsets the positive net pecuniary value differential, $\Delta V_{ic} > 0$, for many students. This is the migration wedge of Section 3; Fact 1 shows it is negative on average and measures the resulting migration penalty.

The remainder of this section establishes three results about this wedge. First, it is large in stated preferences: in a discrete choice experiment that holds pecuniary affordability inside the budget, the average student would forgo a 129% increase in

post-graduation earnings rather than study away from home, a penalty that rises with distance, from 39% for the closest option to 303% for the furthest option, and female students would forgo substantially more than male students. Second, the wedge is anchored in systematic belief differentials between home and away destinations on specific non-pecuniary outcomes (above all safety, parental approval, and expected happiness), which together account for about 61.3% of it. Third, students hold these beliefs in a sparse information environment: most have no source of college advice, few networks in destination cities, and limited knowledge of colleges away from home.

4.1 Fact 1: Migration is costly

To estimate the non-pecuniary cost differential $\Delta N_{i,c}$, I use a discrete choice experiment (DCE) embedded in the baseline survey. Each student is presented with a sequence of pairwise choices between hypothetical college profiles, together with a “No university” opt-out.

Each college is characterized by a major, a measure of educational quality, a destination within Maharashtra, a tuition fee, and monthly living expenses. Major matches the student’s preferred field. Tuition and living costs are held constant within bounds set by the student’s previously stated budget, so profiles are not refused on the basis of affordability. Destination and educational quality are randomly varied. The exercise therefore asks students to trade off destination (*home* or somewhere *away*) against educational quality while holding feasibility fixed. The profiles are not drawn from the actual college landscape: a student may face a top-ranked college at home against a mid-ranked college in Pune even though no such college exists at home. In reality quality and destination are correlated (better colleges are disproportionately in the distant cities), so choices over realistic profiles could not separate a taste for quality from an aversion to migrating; randomizing quality independently of destination breaks that correlation and identifies the migration wedge.

I implement this tradeoff in three quality-numeraire arms. A natural numeraire is expected earnings after graduation, but this may not be the quality signal students actually observe when choosing colleges. I therefore also use placement rates and admission cutoffs, which are more readily available and may guide students’ choices. To reduce cognitive load, each student is randomly assigned to one arm and sees only one quality measure. Before the experiment, students watch a video that explains how to interpret the question and the profiles. Appendix D details the destination-set construction, the attribute levels for each arm, and the budget-anchoring procedure; and shows the screens as students saw them and provides details on the instructions to students. The main

estimates use the earnings arm, pooled with a midline re-administration of the same design to control students that adds the student’s own district capital as a destination (Appendix D); the employment and cutoff arms are reported in Appendix D.1.

I specify utility as linearly additive in destination and the salary numeraire and estimate a conditional logit, following Folke and Rickne (2022):

$$U_{ijt} = \beta_0 \text{NoUniversity}_j + \delta' \mathbf{D}_j + \lambda' \mathbf{S}_j + \varepsilon_{ijt},$$

where U_{ijt} is the latent utility student i assigns to alternative j in choice set t , $\mathbf{D}_j$ collects indicators for the destination of alternative j relative to home (the district capital, Pune, Mumbai, and Nagpur, or a single indicator for any away option), $\mathbf{S}_j$ the salary levels relative to the expected level (-25 , $+25$, and $+50$ percent), and ε_{ijt} is type-I extreme value. Within a choice set the only varying attributes are destination and salary, so δ_d identifies the utility cost of destination d relative to home. The marginal utility of a one-percent salary increase is recovered from the level coefficients,

$$\hat{\beta}_{\% \text{Salary}} \approx \frac{1}{3} \left(-\frac{\hat{\lambda}_{-25}}{25} + \frac{\hat{\lambda}_{+25}}{25} + \frac{\hat{\lambda}_{+50}}{50} \right),$$

and the marginal rate of substitution, the salary increase that compensates a student for studying in d rather than at home, is $\text{MRS}_d \approx -\hat{\delta}_d / \hat{\beta}_{\% \text{Salary}}$; the logit scale cancels in the ratio. Analogous calculations price the penalty in percentage points of employment and in rupees (Appendix D.1).

Table 1 reports the estimates, overall and by gender, and panels (a) and (b) of Figure 3 plot them by destination. The away indicator is large, negative, and significant. The average student would forgo a 129% increase in post-graduation earnings (s.e. 26.3) rather than study away from home, and the penalty rises steeply with distance: 39% for the district capital, 93% for Pune, 196% for Mumbai, and 303% for Nagpur. The ranking is the same when the numeraire is the probability of employment (a 24 percentage-point increase to leave home) or rupees of salary rather than percent of the student’s own expectation (Appendix Tables A6 and A7), and in each baseline arm estimated on its own (Appendix Table A5). The admission-cutoff arm points the same way, but I read it more cautiously. Taken at face value the penalty is 231 program ranks, but students do not treat a higher cutoff as a desirable quality signal: they choose more-selective programs less. I therefore do not interpret the cutoff-arm willingness to pay as a quality tradeoff.⁷ Female students would give up substantially more than male students: a 227%

⁷Students’ aversion to higher cutoffs could have several sources: they may not have fully internalized the instruction to assume they have scored high enough marks to be admitted to both profiles and may instead be reacting to the perceived difficulty of

versus 72% earnings increase to leave home, a gap that widens with distance (panel (b) of Figure 3). The gap comes from the numeraire rather than from the move itself. The utility cost of leaving home is similar for men and women (-0.553 versus -0.614), while men's salary coefficient is 2.8 times women's, so the same utility cost translates into a larger share of expected earnings for women; priced in the probability of employment, which men and women value alike, the gender gap all but disappears except for Mumbai (Appendix Table A6), and priced in rupees the leave-home ask is about the same for both (Appendix Table A7). A complementary incentivized college-rating exercise that does not anchor pecuniary attributes to expectations, implies a willingness to pay to stay home of the same order as the DCE, with the same steep gender gap (Appendix F).⁸ Fact 2 returns to this gap.

These stated-preference estimates establish that the migration penalty is large, and that women face a substantially higher penalty than men. Large as it is, this penalty is still only about half the disutility of forgoing college altogether: in the baseline earnings arm, the median student treats not attending university as roughly 2.0 times as costly as a move away from home (a college at home is worth about 230% of expected earnings over no college at all), so the ranking is home, then away, then no university.

4.2 Fact 2: Beliefs about non-pecuniary outcomes drive much of the penalty

Fact 1 shows the migration wedge is large. I next ask what it is anchored in. For each destination a student considers, the survey asks how they picture life there along a set of non-pecuniary dimensions, which I group into three families. *Emotional well-being* is how happy they expect to be and how likely they are to struggle with their mental health. *Social ties* covers how often they would see family and friends back home, how many close friends they would make, whether people back home would think worse of them for leaving, and whether their parents would approve. *Practical and lived experience* covers how safe they would feel and how hard the move itself would be. Each item is elicited on a 0–10 scale, a probability, or a count, and I recode every item so that higher

getting in; a higher cutoff may signal selectivity without a commensurate labor-market return; a more selective program may be seen as harder to study at; or applying to distant, more competitive programs may simply be more difficult. I cannot separate these. Only the first would threaten the identification of the wedge, so as a robustness check I re-estimate the away penalty controlling for each student's perceived probability of admission at every destination. Even in that case, that is, even if students disregarded the assume-admission instruction entirely, the penalty remains large and highly significant, at 65% of earnings. The same aversion to selectivity appears in the complementary college rating exercise, where the cutoff was drawn independently of the student's own expected score (Appendix F).

⁸In the rating exercise, students score hypothetical college profiles (0–10) whose destination varies alongside fees, earnings, placement, and admission difficulty (Table A15). Ratings respond to fees but only weakly to the earnings attribute, so I identify the away penalty from the fee students would forgo rather than from earnings: a college away from home is rated about a point lower, an implicit penalty of roughly Rs 86,000 a year in tuition. Expressed against students' own expected earnings, that is about 86% of a year's median expected earnings overall, of the same order as the baseline DCE away penalty of 126% (Appendix Table A5), and 51% for men versus 133% for women, echoing the DCE's gender gap.

values are always better (safer, happier, more approving, and so on). I then standardize each by its home mean and home standard deviation, so a gap reads as how much worse students expect a destination to be than their own district. Appendix E gives the exact question wording and response scales.

Across all eight outcomes, students expect substantially worse experiences away from home than at home (Figure 4, Table A13). The gaps are largest on safety (-0.72 home standard deviations), parental approval (-0.62), and expected happiness (-0.53); the remaining gaps are smaller but all significant at the 1 % level. For women the gaps are wider in parental approval, safety, and reputation back home (Appendix Figure A6). Similar constraints are documented elsewhere in the literature: women in Delhi choose lower-ranked colleges to travel by safer routes (Borker, 2021), women’s enrollment in skills training in rural Pakistan drops sharply once the center lies beyond their own village (Cheema et al., 2025), and workers forgo pay to avoid workplaces with a risk of sexual harassment, a risk borne disproportionately by women (Folke and Rickne, 2022).⁹

If these beliefs are part of the non-pecuniary value of destinations, controlling for them should shrink the migration penalty in the DCE. Panel (c) of Figure 3 re-estimates the destination coefficients with each student’s own beliefs about each destination added as controls, pooling the three baseline arms with the midline choice sets of control students and holding the sample fixed (Appendix D). The beliefs are jointly significant (Wald $p = 0.000$) and absorb most of the penalty. Measured by the MRS to stay home in the earnings-arm model, the leave-home penalty falls from 129 % to 42 % of post-graduation earnings, a reduction of 67.3 %.¹⁰ Figure 5 summarizes the results of this exercise by belief family: in the linear-probability version of the pooled model, the belief families jointly account for 61.3 % of the leave-home coefficient, with the practical items accounting for the largest share (34.0 %), emotional items accounting for 13.3 %, and social items accounting for 14.1 %.¹¹ Item by item (Table A9), moving difficulty, happiness, safety, and parental approval carry the largest shares; the remaining items contribute little.

The gender columns of Table A9 revisit Fact 1’s gender gap. Beliefs absorb more of

⁹A broader literature documents gendered mobility constraints. Providing bicycles to girls in Bihar narrowed the gender gap in secondary school enrollment by making the commute feasible (Muralidharan and Prakash, 2017); women in Lahore apply for jobs at more than double the rate when offered women-only transport (Field and Vyborny, 2022); and women in Rio de Janeiro pay a premium to ride women-only train cars and avoid harassment (Kondylis et al., 2025). Jayachandran (2015) reviews the norms behind these constraints.

¹⁰Moving difficulty is asked only about away destinations and relative to home: the home option takes the value “not challenging at all,” which 12 % of away answers also select. The home value is definitional, so part of this item’s contribution rests on the anchoring.

¹¹The decomposition is the pooled specification of Appendix Table A9 ($N = 1682$ students); the whiskers in Figure 5 are 95 % cluster-bootstrap confidence intervals on each family’s share (1000 replications, resampling students, stratified by gender). The share depends on the sample: it is lowest when the earnings arm is used alone and highest in the cutoffs arm and at midline (Appendix Table A8).

women's penalty (74.1 %) than men's (41.6 %), and the decomposition locates the difference in the practical items: safety and moving difficulty account for 46.5 % of women's penalty against 13.7 % of men's, while the emotional and social families contribute similar shares for both. The same pattern holds in the earnings-arm model: with beliefs held fixed, women's leave-home MRS falls by 98.3 % and men's by 11.7 %. Conditional on what they believe about the practicalities of moving, women are no more home-attached than men.

Perceived non-pecuniary costs of migration are therefore large, systematic, and account for a substantial share of the migration penalty, but they are anchored in students' subjective beliefs about destinations that, as the next subsection shows, students have limited means of verifying.

4.3 Fact 3: Students have limited information and few network connections

Facts 1 and 2 establish that the migration penalty is large and anchored in students' beliefs about non-pecuniary outcomes at the destination. Whether those beliefs are accurate, and whether they can be changed, depends on what students can observe about destinations and from whom. If students have little direct information about destinations, their beliefs may be inaccurate, and peers may offer a way for them to correct beliefs. Alternatively, even if the beliefs are accurate on average, access to someone with destination experience may directly reduce the underlying cost of migration by making the destination easier to understand and navigate. This subsection documents that the broader information environment in which students form destination-specific beliefs is sparse: most students report no source of college advice, few have networks in destination cities, and students have low absolute levels of knowledge about both the college landscape and the destinations themselves, with the gaps widest for the distant cities.

About 29 % of students report no source of college advice at all, and 30 % have not yet considered any specific colleges at baseline. Only 4 of the 52 schools in the sample employ a dedicated college counselor, so only about 12 % of students have institutional support inside the school. Even at home, students can on average name only 1.9 colleges in their preferred field (1.5 in the district capital for students living outside it), and when asked about a specific, randomly chosen college there, only 43 % have heard of it (Figure 6). Appendix Figure A7 shows the broader picture: most students report low confidence in navigating the admissions process.

The knowledge gaps are larger for the distant cities, where the decision to enroll requires the most information. The count of colleges students know within their preferred field falls with distance, despite the fact that selective colleges are disproportionately lo-

cated in these locations. Students' broader knowledge of educational hubs in their state is also limited. Asked how well they know each of the cities the survey randomly assigns them, 60 % of the responses is "I don't know anything about it"; 41 % of students do not have any relatives or friends in any of the destinations asked.

Together, these three facts describe a setting in which students face a large migration penalty, anchored in beliefs about non-pecuniary outcomes they cannot easily verify. The randomized intervention described in Section 5 targets exactly this gap by introducing students to peers who have already made such a move.

5 Experimental design and data

The experiment evaluates Campus Trail, a phone-based peer-mentorship program designed to expand students' access to migrant peers before college application deadlines. Eligible students are in their final year of secondary school and plan to continue to higher education. Treated students are offered a match with a current university student in Pune or Mumbai who migrated for college and can discuss college choice, applications, and the practical experience of studying away from home. The program is therefore designed to target the information and network gaps documented in Section 4.

Sampling I offer the program to every secondary school in the two study districts. Working from lists provided by the District Education Officers, the field team first reaches each headmaster by phone to introduce the program and then visits the schools whose headmasters consent. During a visit, the team presents Campus Trail and invites interested students to sign up through a short tablet survey that doubles as the program application; when more students sign up than can be surveyed in a day, the team selects among them at random, stratifying by stream and gender. To incentivize truthful answers, students and parents are told that their responses may be used to match students to mentors, so that more accurate reporting yields a better match if the student is selected (Low, 2024; Kessler et al., 2019).

Randomization The randomization sample consists of 1,682 students across 52 schools who meet pre-specified interest and attention criteria (Appendix G). I randomly assign half of the shortlisted applicants to Campus Trail and half to control, 826 and 856 students respectively. The two arms are well balanced on baseline characteristics: no individual difference is significant, and a joint orthogonality test does not reject (Appendix Table A17).

Intervention Mentors are current or recent university students in Pune and Mumbai. Appendix B describes the applicant pool. They join as paid part-time interns after a short training on program rules, safeguarding, and the calling tools. Each treated student is matched to a mentor of the same gender and, where possible, one who studied the student’s intended major or destination or shares their academic stream. The pair then meets by phone. Conversations cover college and major choice, the practicalities of moving (housing, safety, finding one’s way around an unfamiliar city), and encouragement drawn from the mentor’s own experience. The mentor-applicant data suggest that these peers brought a rich stock of concrete, destination-specific knowledge as well as reassurance that could operate through the three channels outlined in Section 3: Mentors are prepared to offer information that could lead to belief updating by supplying concrete destination-specific information on budgets, accommodation, transport, safety, scholarships, and academic returns. They could reduce effective migration costs by lowering search and coordination costs: identifying hostels or PGs, explaining procedures, suggesting routes, naming financial-support schemes, and connecting students to seniors or other contacts. Finally, they could affect how students weight anticipated costs by acknowledging that moving may involve homesickness, loneliness, safety concerns, or parental resistance, while making compensating benefits salient: new friendships, broader peer networks, independence, exposure, and career opportunities. Appendix B documents these examples in detail. Mentors’ commitment is capped at three hours per mentee, for which mentors receive a 200 INR contact bonus and a further 800 INR on completion (about 10 USD in all). Students who participated in conversations with their mentors reported positive impressions: among treated students who had spoken with their mentor by the midline survey, 91% agreed the mentor gave them useful information and advice, 90% felt encouraged after their sessions, and 85% said they trusted their mentor. Their conversations spanned the intended range of topics, most often academics (46%) and career prospects such as internships and placements (43%), but also the cost of living (24%), safety (22%), housing and daily logistics (20%), and navigating family wishes about studying away from home (19%).

Take-up Calls are conducted through a SurveyCTO integration that connects the pair keeping mentors’ contact details anonymous and logs each attempt’s duration and outcome. I define take-up as at least fifteen minutes of conversation verified in these logs. The first conversation is scheduled by the field team, and enumerators attempted contact with every student assigned to the mentorship program. 92% connected with their mentor at least once, and 31% went on to have a recorded conversation longer than fifteen

minutes, averaging 41 minutes. These figures understate true engagement, because mentors sometimes called from personal phones or failed to record a call. Supplementing the logs with mentor-reported durations where possible implies that 45% of students had a call longer than fifteen minutes, with 40 minutes of engagement on average. Among treated students who did not take up, enumerators logged an average of six scheduling attempts and mentors placed an average of five to six calls. Engagement also accumulated gradually over the intervention, so the dose in place at the midline survey was smaller than these end-of-intervention figures. Among treated students interviewed at midline, one in four (25%) had completed a recorded fifteen-minute conversation by their interview date. That is about three-quarters of the recorded take-up this group reached by the end of the intervention, so the short-run midline results reflect a partial dose. Take-up follows students' own plans: a one-standard-deviation increase in a student's baseline likelihood of applying to colleges away from home raises it by about seven percentage points, and take-up rises with academic preparation and self-efficacy but not with household resources (Appendix Table A18).¹²

5.1 Data

The data come from three survey waves, together with the mentorship call records just described. The baseline survey, administered on tablets in school in August and September 2025, precedes the intervention and collects household and demographic characteristics, academic background, the discrete choice experiment and college rating exercise, destination-specific beliefs, measures of information and search, and trait batteries covering self-efficacy, gender attitudes, aspirations, and social desirability (the Marlowe–Crowne scale). Alongside the student baseline, I survey one parent or guardian per student by phone, reaching a randomly designated mother or father. 47% (792 of 1,682) of students have a parent interview, with parents of more engaged households selecting into the survey, but it does not differ across treatment arms (Appendix H). Two administrative sources supplement the surveys: school characteristics used to benchmark the study schools against others in the two districts come from the Unified District Information System for Education (UDISE+), and the higher-education supply figures in Section 2 come from public entrance-exam and seat-allocation records.

The midline survey, administered in schools between November 2025 and January

¹²These associations are individually significant but collectively weak. A few baseline traits predict take-up (students who aspire to engineering or postgraduate study, or who are more willing to take risks, engage somewhat more), but across the full baseline covariate set take-up is essentially unpredictable out of sample (five-fold cross-validated $R^2 \approx 0$; random-forest out-of-bag $R^2 \approx 0.02$), as are total call length and mentors' engagement ratings, and engagement varies little across mentors (intra-class correlation ≈ 0.03). Who engages with their mentors, and how much, is close to unpredictable from what is observed at baseline.

2026, before entrance exams and admissions, re-surveyed 1,204 students (72% of the sample). It measures the hypothetical-migration outcomes and the belief, information, and trait mechanisms, and re-administers the discrete-choice experiment and college-rating exercise, the source of the structural willingness-to-pay estimates reported below; at this stage I elicit beliefs about home and the district capital rather than the district as a whole (Appendix D). Because it was conducted in person during school visits, the midline involved no tracking: enumerators surveyed the students who were present. The endline survey, conducted by phone from July 2026 as admissions conclude, measures realized migration and enrolment, with a short proxy interview with a parent, friend, or teacher standing in (flagged and checked for robustness) when a student cannot be reached directly.

1,204 students (72%) completed the midline survey and 1,143 (68%) completed a direct endline interview, rising to 1,418 (84%) once endline proxy interviews are included. Midline response is selective on which students attended school on the survey day but does not differ by treatment arm; endline response, reached by phone, shows mild differential attrition among treated students with the strongest baseline migration intentions. Appendix Table A20 reports response rates by arm and wave, Appendix Table A22 presents tests for selective and differential attrition, and the endline tracking protocol and attrition bounds are described in Appendix I.

6 Empirical strategy

The analysis follows the pre-analysis plans filed before the midline and endline data were accessed. This section describes the primary outcome, the estimating equation, and the analysis samples.

Primary outcome The primary outcome is an inverse-covariance-weighted index of standardized components (Anderson, 2008). Each component is standardized to the control-group mean and standard deviation; the weights are the row sums of the inverted control-group covariance matrix; and the index is re-standardized to mean zero and unit standard deviation in the control group. For respondents missing a component, the index aggregates the observed components only, with weights recomputed from the inverse of the corresponding covariance submatrix. The index combines six indicators measured at endline: having moved away for study, being enrolled or planning to enroll in higher education, migrating for education, applying to college, applying to a college away from home, and seeking financing for their education. Each indicator counts a real-

ized action or, for students who have not yet acted, a planned application or move within the following two months. I report each outcome for any move away from home, which includes the district capital; the out-of-district definition is in Appendix Table A25.¹³ Exact component definitions are specified in the endline pre-analysis plan.

Estimating equation For student i I estimate the intent-to-treat (ITT) effect by OLS:

$$Y_i = \alpha + \beta \text{CampusTrail}_i + \tilde{X}_i' \gamma + (\text{CampusTrail}_i \times \tilde{X}_i)' \delta + \varepsilon_i, \quad (1)$$

where CampusTrail_i indicates assignment to mentorship and β is the ITT. Following Lin (2013), the covariates are mean-centered ($\tilde{X}_i$) and interacted with treatment. $\tilde{X}_i$ contains the baseline value of the outcome where available and additional controls chosen by post-double-selection lasso (Belloni et al., 2014) from a pre-specified candidate set (baseline and expected grades, assets, indices of agency, knowledge, and effort, time, risk, and competitiveness preferences, number of close friends, educational expectations, and baseline values of all outcomes), with missing baseline values set to zero and flagged by an indicator. I report heteroskedasticity-robust standard errors for student-level outcomes and cluster by student for the destination-stacked belief specifications, which carry multiple observations per student.

Heterogeneous effects I report the pre-specified heterogeneity by gender, and an exploratory split by whether the student lives in a district capital.¹⁴

Analysis samples and attrition The endline sample comprises students reached after admissions conclude. Indices are computed on complete cases. Endline response is mildly differential: treated students who intended to stay home at baseline are harder to reach, which may introduce bias to my estimates; imputation exercises reported in Appendix I move the index by about one hundredth of a standard deviation.

Transformations Continuous monetary variables are winsorized at the 99th percentile within treatment arm and district, expressed in thousands of rupees, and transformed by the inverse hyperbolic sine; I drop variables that take a single value for at least 95% of the relevant sample.

¹³Home is the student's own sub-district, as defined in Section 2. A move away from home denotes moving beyond this local area, while an out-of-district move additionally crosses the broader district boundary.

¹⁴The pre-analysis plan also specifies heterogeneity by academic stream and by the existence of baseline networks away from home; these are not reported in this draft.

7 Effects of mentorship

7.1 Educational migration

Table 2 reports intention-to-treat effects on the pre-specified migration index and its six components. Assignment to mentorship raises the index by 0.16 standard deviations. On the extensive margin, treated students are three percentage points more likely to have applied for further study and five points more likely to be enrolled or planning to enroll, against control means of 91 and 79 percent. On the intensive margin, where they study, every estimate is positive but imprecise: four percentage points on applying to a college away from home and three on moving away. The estimate is stable across constructions of the index, from 0.13 to 0.20 standard deviations across the post-double-selection specification, direct interviews only, the expanded sample, and realized outcomes only (Appendix Table A24), and it survives the attrition adjustments of the pre-analysis plan: 0.12 under intensive-tracking weights, and between 0.10 and 0.15 under adversarial sensitivity bounds (Appendix Table A23).

7.2 Where the induced moves go

Table 3 breaks the applications and moves down by destination. The additional applications increase only for the district capital: five percentage points on a base of 11 percent, with no change for home, elsewhere in the home district, Pune and Mumbai, or other districts. Moves follow the same pattern, migration to the district capital increases by two percentage points, with no increases to the cities where mentors' study. The treatment effects are concentrated entirely on students who live outside district capitals at baseline, which consist of approximately three quarters of the sample: their index rises by 0.21 standard deviations, while students already living in one show no effect (Appendix Table A27). For these students, a move to the district capital is the option that carries the smallest migration penalty (Section 4).

7.3 Men and women

Table 4 splits the estimates by gender, on direct student interviews. Men's applications away from home increases by ten percentage points, and migration increases by four percentage points. Women show no effect on any migration outcome.

7.4 Beliefs about destinations

Table 5 reports effects on endline beliefs about the district capital and about Pune, relative to home. On average the program does not improve how students picture life away from home: non-pecuniary beliefs are flat at the district capital and, if anything, lower at Pune, and beliefs about returns do not move. The one belief that moves is the cost of the district capital, which treated students expect to be 0.42 standard deviations cheaper relative to home. Outside the district capitals the pattern mirrors the actions: men revise the cost of the district capital down by 0.74 standard deviations, and women revise their non-pecuniary view of Pune down by 0.47. Contact with a migrant peer made the near destination look cheaper to the students who then applied there, and made the far destination look worse to those who did not move. No other pre-specified channel responds: knowledge of destinations, information about admissions, search effort, support, and traits are unchanged in both waves (not reported).

8 Conclusion

This paper studies the reluctance to migrate at a critical point in life: the choice of where to study for college. Among final-year secondary students in Maharashtra, the perceived penalty of migration is large: the average student would forgo a 129% increase in expected post-graduation earnings to avoid studying away from home, and beliefs about the difficulty of the move, safety, parental approval, and happiness accounts for a large share of the penalty. Despite expecting higher returns elsewhere, students plan to study at home, and this penalty rationalizes the choice. A light-touch intervention consisting of conversations with a college student who is an educational migrant changes choices, but only at the margin where the perceived cost was small to begin with. It does not change the penalty attached to the distant cities, where the pecuniary gains concentrate, and if anything beliefs about life there move in the wrong direction. If those beliefs are broadly correct, and worsening on contact with someone who made the move suggests they are, encouraging educational migration at scale means addressing the costs themselves, safe housing and safe travel among them, rather than correcting misperceptions.

References

- Acton, R., K. E. Cortes, L. Miller, and C. Morales (2025). Distance to degrees: How college proximity shapes students' enrollment choices and attainment across race-ethnicity and socioeconomic status. *Economics of Education Review*, 102724.
- Alba-Vivar, F. M. (2024). Opportunity bound: Transport and access to college in a megacity.
- Alfonsi, L., M. Namubiru, and S. Spaziani (2022). Meet your future: Experimental evidence on the labor market effects of mentors. Working Paper 87, G²LM|LIC.
- Anderson, M. L. (2008). Multiple inference and gender differences in the effects of early intervention: A reevaluation of the abecedarian, perry preschool, and early training projects. *Journal of the American Statistical Association* 103(484), 1481–1495.
- Aydemir, A. B., M. G. Kirdar, and H. Torun (2022). The effect of education on internal migration of young men and women: Incidence, timing, and type of migration. *Labour Economics* 74, 102098.
- Baláž, V., A. M. Williams, and E. Fifeková (2016). Migration decision making as complex choice: Eliciting decision weights under conditions of imperfect and complex information through experimental methods. *Population, Space and Place* 22(1), 36–53.
- Barnett-Howell, Z., T. Baseler, T. Ginn, and S. Gordeev (2025). Reaching the novice or nudging the expert? networks, information, and the experimental returns to migration. *Unpublished manuscript*.
- Batista, C., D. M. Costa, P. Freitas, G. Lima, and A. B. Reis (2025). What matters for the decision to study abroad? A lab-in-the-field experiment in Cape Verde. *Journal of Development Economics* 173, 103401.
- Beaman, L., E. Duflo, R. Pande, and P. Topalova (2012). Female leadership raises aspirations and educational attainment for girls: A policy experiment in india. *Science* 335(6068), 582–586.
- Beber, B., C. Ebert, Z. Riaz, and J. Singh (2026). Improving migration outcomes: A mentoring experiment and its network effects in Senegal. Work in progress.
- Belloni, A., V. Chernozhukov, and C. Hansen (2014). High-dimensional methods and inference on structural and treatment effects. *Journal of Economic Perspectives* 28(2), 29–50.
- Bettinger, E. P., B. T. Long, P. Oreopoulos, and L. Sanbonmatsu (2012). The role of application assistance and information in college decisions: Results from the H&R Block FAFSA experiment. *The Quarterly Journal of Economics* 127(3), 1205–1242.

- Blumenstock, J. E., G. Chi, and X. Tan (2025, January). Migration and the value of social networks. *The Review of Economic Studies* 92(1), 97–128.
- Boneva, T. and C. Rauh (2017). Socio-economic gaps in university enrollment: The role of perceived pecuniary and non-pecuniary returns. CESifo Working Paper 6756, CESifo.
- Borker, G. (2021). Safety first: Perceived risk of street harassment and educational choices of women.
- Breda, T., J. Grenet, M. Monnet, and C. Van Effenterre (2023). How effective are female role models in steering girls towards STEM? evidence from french high schools. *The Economic Journal* 133(653), 1773–1809.
- Bryan, G. and M. Morten (2019). The aggregate productivity effects of internal migration: Evidence from Indonesia. *Journal of Political Economy* 127(5), 2229–2268.
- Caner, A., M. Derebasoglu, and C. Okten (2024). Attainment and gender equality in higher education: Evidence from a large-scale expansion. *Journal of Human Capital* 18(3), 469–530.
- Card, D. (1995). Using geographic variation in college proximity to estimate the return to schooling. In L. N. Christofides, E. K. Grant, and R. Swidinsky (Eds.), *Aspects of Labour Market Behaviour: Essays in Honour of John Vanderkamp*, pp. 201–222. Toronto: University of Toronto Press.
- Carrell, S. E. and B. Sacerdote (2017). Why do college-going interventions work? *American Economic Journal: Applied Economics* 9(3), 124–151.
- Castleman, B. L. and L. C. Page (2015). Summer nudging: Can personalized text messages and peer mentor outreach increase college going among low-income high school graduates? *Journal of Economic Behavior & Organization* 115, 144–160.
- Cheema, A., A. I. Khwaja, M. F. Naseer, and J. N. Shapiro (2025). Glass walls: Experimental evidence on constraints faced by women in accessing valuable skilling opportunities. *Journal of Political Economy*. Forthcoming.
- Clayton, K., Y. Horiuchi, A. R. Kaufman, G. King, and M. Komisarchik (2023). Correcting measurement error bias in conjoint survey experiments. Working paper.
- Clerici, C., I. Bianchi, and D. Biesalski (2024). Information, expectations and preferences: Occupational choices of young adults in Uganda. Working paper, Stockholm School of Economics.
- Conzelmann, J. G., S. W. Hemelt, B. J. Hershbein, S. Martin, A. Simon, and K. M. Stange (2025). Grads on the go: Measuring college-specific labor markets for graduates. *Journal of Policy Analysis and Management* 44(3), 741–763.
- Coppock, A., A. S. Gerber, D. P. Green, and H. L. Kern (2017). Combining double


- sampling and bounds to address nonignorable missing outcomes in randomized experiments. *Political Analysis* 25(2), 188–206.
- Delavande, A. and B. Zafar (2024). University choice: The role of expected earnings, nonpecuniary outcomes, and financial constraints. *Journal of Political Economy*. Forthcoming.
- Duflo, E. (2001). Schooling and labor market consequences of school construction in Indonesia: Evidence from an unusual policy experiment. *American Economic Review* 91(4), 795–813.
- Elsayed, A. and A. Shirshikova (2023). The women-empowering effect of higher education. *Journal of Development Economics* 163, 103101.
- Fabre, A. (2024). The geography of higher education and spatial inequalities.
- Farre, L. and F. Ortega (2021). Family ties, geographic mobility and the gender gap in academic aspirations.
- Field, E. and K. Vyborny (2022). Women’s mobility and labor supply: Experimental evidence from Pakistan.
- Folke, O. and J. Rickne (2022). Sexual harassment and gender inequality in the labor market. *The Quarterly Journal of Economics* 137(4), 2163–2212.
- Fu, C., J. Guo, A. J. Smith, and A. T. Sorensen (2022). Students’ heterogeneous preferences and the uneven spatial distribution of colleges. *Journal of Monetary Economics* 129, 49–64.
- Gelbach, J. B. (2016). When do covariates matter? and which ones, and how much? *Journal of Labor Economics* 34(2), 509–543.
- Ghose, D. (2026). Trade, internal migration, and human capital: Who gains from India’s IT boom? Working paper, World Bank.
- Giustinelli, P. (2023). Expectations in education. In *Handbook of Economic Expectations*, pp. 193–224. Academic Press.
- Gollin, D., D. Lagakos, and M. E. Waugh (2014). The agricultural productivity gap. *The Quarterly Journal of Economics* 129(2), 939–993.
- Green, A. (2024, November). Networks and geographic mobility: Evidence from World War II navy ships. Job Market Paper, Princeton University. Available at https://allie-e-green.github.io/public/WW2_AG_Migration_jump.pdf.
- Han, Y., D. Kübler, Y. Liu, X. Zheng, and Y. Zong (2025, June). The gender gap in gender-blind university admissions. Working paper.
- Hole, A. R. (2015). DCREATE: Stata module to create efficient designs for discrete choice experiments. Statistical Software Components S458059, Boston College Department of Economics. Revised 2017.

- Hoxby, C. M. and S. Turner (2013). Expanding college opportunities for high-achieving, low income students. Discussion Paper 12-014, Stanford Institute for Economic Policy Research.
- Ishimaru, S. (2025). Geographic mobility of youth and spatial gaps in local college and labor market opportunities. *Journal of Labor Economics* 43(4), 1251–1294.
- Jayachandran, S. (2015). The roots of gender inequality in developing countries. *Annual Review of Economics* 7(1), 63–88.
- Jensen, R. (2010). The (perceived) returns to education and the demand for schooling. *Quarterly Journal of Economics* 125(2), 515–548.
- Kapoor, R. and S. Gade (2024). Together in search: Experimental evidence from coordinating travel among women job-seekers in urban India.
- Kennan, J. and J. R. Walker (2011). The effect of expected income on individual migration decisions. *Econometrica* 79(1), 211–251.
- Kessler, J. B., C. Low, and C. D. Sullivan (2019). Incentivized resume rating: Eliciting employer preferences without deception. *American Economic Review* 109(11), 3713–3744.
- Koenen, M. and D. Johnston (2024, December). Social ties and residential choice: Micro evidence and equilibrium implications. Working Paper.
- Kondylis, F., A. Legovini, K. Vyborny, A. Zwager, and L. C. d. Andrade (2025). Demand for “safe space”: Avoiding harassment and complying with norms. *Journal of Development Economics* 174, 103392.
- Kone, Z. L., M. Y. Liu, A. Mattoo, Ç. Özden, and S. Sharma (2018). Internal borders and migration in India. *Journal of Economic Geography* 18(4), 729–759.
- Koşar, G., T. Ransom, and W. Van der Klaauw (2022). Understanding migration aversion using elicited counterfactual choice probabilities. *Journal of Econometrics* 231(1), 123–147.
- Lagakos, D. (2020). Urban-rural gaps in the developing world: Does internal migration offer opportunities? *Journal of Economic Perspectives* 34(3), 174–192.
- Lagakos, D., A. M. Mobarak, and M. E. Waugh (2023). The welfare effects of encouraging rural-urban migration. *Econometrica* 91(3), 803–837.
- Le Barbanchon, T., R. Rathelot, and A. Roulet (2021). Gender differences in job search: Trading off commute against wage. *Quarterly Journal of Economics* 136(1), 381–426.
- Lee, D. S. (2009). Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *The Review of Economic Studies* 76(3), 1071–1102.
- Lin, W. (2013). Agnostic notes on regression adjustments to experimental data: Reexamining freedman’s critique. *The Annals of Applied Statistics* 7(1), 295–318.

- Low, C. (2024). Pricing the biological clock: The marriage market costs of aging to women. *Journal of Labor Economics* 42(2), 395–426.
- Mangal, K. (2024). How much are government jobs in developing countries worth? *The World Bank Economic Review* 38(3), 443–465.
- Manski, C. F. (2004). Measuring expectations. *Econometrica* 72(5), 1329–1376.
- McKenzie, D. (2022). *Fears and tears: Should more people be moving within and from developing countries, and what stops this movement?* The World Bank.
- Mesplé-Somps, S. and B. Nilsson (2023). Role models, aspirations and desire to migrate. *Journal of Economic Behavior & Organization* 212, 819–839.
- Ministry of Human Resource Development (2013). All india survey on higher education 2011–12. Technical report, Department of Higher Education, Government of India, New Delhi.
- Munshi, K. and M. Rosenzweig (2016, January). Networks and Misallocation: Insurance, Migration, and the Rural-Urban Wage Gap. *American Economic Review* 106(1), 46–98.
- Muralidharan, K. and N. Prakash (2017). Cycling to school: Increasing secondary school enrollment for girls in India. *American Economic Journal: Applied Economics* 9(3), 321–350.
- Riley, E. (2024). Role models in movies: The impact of Queen of Katwe on students' educational attainment. *The Review of Economics and Statistics* 106(2), 334–351.
- Sahai, H. and M. Bailey (2022). Social networks and spatial mobility: Evidence from facebook in india. *arXiv preprint arXiv:2203.05595*.
- Sekhri, S., M. A. Hossain, and P. Khosla (2024). Access to colleges, human capital, and empowerment of women. *Journal of Human Resources* 59(2), 502–544.
- Sjaastad, L. A. (1962). The costs and returns of human migration. *Journal of Political Economy* 70(5), 80–93.
- Wiswall, M. and B. Zafar (2015). Determinants of college major choice: Identification using an information experiment. *The Review of Economic Studies* 82(2), 791–824.


Figures

Figure 1: Study sites: Ahilyanagar and Sangli districts.




Notes: Maharashtra, with the two study districts shaded. Dots are the 52 study schools.

Figure 2: Access to engineering and business-program seats from Sangamner, by applicant percentile and geography.



Notes: The panels report the number of 2025–26 CAP seats accessible to a hypothetical applicant based in Sangamner taluka at each entrance-exam percentile. The applicant is male, general category, non-EWS, and has no TFWS, disability, defence, or orphan entitlement; Sangamner is assigned to the Savitribai Phule Pune University home-university region. Business programs comprise BBA, BMS, and BBM. The three geographic categories are mutually exclusive: home (Sangamner taluka), other talukas in Ahilyanagar district, and all other Maharashtra districts. A seat is counted when the applicant satisfies its gender, category, special-entitlement, and home-/other-than-home-university rules and meets the terminal CAP-IV cutoff; terminal-round vacancies are assigned a cutoff of zero, while tranches without a comparable cutoff are excluded. The vertical axis is logarithmic. The curves cover 96.3 % of modeled engineering seats and 90.6 % of modeled business-program seats that can be assigned to a taluka from institute-address evidence. Counts describe the size of the feasible seat set, not admission probabilities. Source: Maharashtra State Common Entrance Test Cell admissions records and the institute–taluka register constructed from published addresses.

Figure 3: The migration penalty by destination: MRS, MRS by gender, and destination coefficients with and without belief controls.




Notes: Panels (a) and (b): the marginal rate of substitution by destination from the model of Table 1 (baseline earnings arm and midline control choice sets, one indicator per destination relative to the student's own home), the destination coefficient divided by the salary slope, in percent of expected salary; panel (a) pools all students and panel (b) re-estimates the model on men (teal) and women (violet). Panel (c): the destination coefficients themselves, in utility units, from the pooled model that stacks all three baseline numeraire arms with the midline control choice sets (arm-specific attribute blocks and opt-out constants, destination coefficients common across arms and waves), estimated on the choice sets with complete belief data without controls (gray) and adding the eight destination-specific beliefs of each wave (blue); Appendix Figure A5 shows the salary-level coefficients of the same two models and Appendix Table A8 the corresponding penalty reductions by sample. Whiskers are 95 % confidence intervals with standard errors clustered by student, delta-method for the ratios; stars follow the p -value of the destination coefficient. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.


Figure 4: Beliefs about non-pecuniary outcomes by destination

Emotional

Would feel happy there (0–10)




Mental wellbeing (% reverse coded)




Social


Parents would approve (0–10)




Nr. of close friends



Days per month seeing family and friends




People back home would not think worse (%)




Practical

Would feel safe there (0–10)




Would be easy to move there (3–0, reverse coded)




Notes: Each panel plots the mean belief at each destination, in the item's own units, along the distance ladder: the student's own town, the capital of the student's own district (students who live in the capital contribute to the home point only), Pune, Mumbai, and Nagpur, with approximate travel times from the study districts. Means pool the baseline city module (all students: the district capital and two of Pune, Mumbai, and Nagpur) and the midline control group (the own town and two of the district capital, Pune, and Mumbai), adjusted for the survey wave. Mental well-being is the reverse-coded probability of mental-health difficulties; reputation is the probability that people back home would not think worse of the student; moving ease reverses the four-point difficulty item, and its home value is definitional ("not challenging at all") rather than asked. Appendix Table A13 reports the away-home gaps in standard-deviation units; Appendix Figure A6 shows the same panels by gender.

Figure 5: Decomposition of the pooled leave-home penalty (β_{Away}) into belief families.



Notes: β_{Away} is the coefficient on a single indicator for any away option in the linear-probability version of the pooled model of panel (c) of Figure 3 (all baseline arms and midline control), with choice-set fixed effects, on the complete-belief choice sets; it is set to 100 %. Each bar removes the share attributed to one belief family in the decomposition of Table A9 (coefficients in the model with beliefs times within-choice-set belief gaps, in the spirit of Gelbach (2016)); the last bar is the share of the penalty that remains once the beliefs are held fixed.

Figure 6: Baseline knowledge of destinations and of colleges, by destination



Notes: Each panel plots the mean of one item along the distance ladder of Figure 4: the student's own home, the capital of the student's own district, Pune, Mumbai, and Nagpur, pooling the baseline city module (all students) and the midline control group, adjusted for the survey wave. Baseline students living outside a district capital were asked about their own district as a whole; their answers are placed at the district-capital level, and students living in the capital contribute to the home point. Top row: "How much do you know about studying and student life in [city]?" (1 = I don't know anything about it, 4 = I know it really well) and the share with any relatives, family friends, or friends there; both are asked only about away destinations, so the home value is definitional and drawn as an open marker. Bottom row: the number of colleges the student can name in their preferred field there (0 = none, 11 = more than 10) and the share who have heard of a college drawn at random from that destination. Nagpur is asked only at baseline.

Tables

Table 1: The migration penalty: pooled destination-choice estimates

	All	Men	Women
<i>Panel A. Utility coefficients</i>			
Away from home	-0.587*** (0.054)	-0.553*** (0.085)	-0.614*** (0.070)
Salary (per 1% of expected)	0.0046 (0.0008)	0.0077 (0.0013)	0.0027 (0.0010)
<i>Panel B. MRS: % increase in expected salary to migrate</i>			
Away from home	129*** (26.3)	72*** (16.0)	227*** (92.8)
District capital	39** (19.1)	28* (15.9)	60 (48.1)
Pune	93*** (21.9)	63*** (16.0)	149*** (67.0)
Mumbai	196*** (39.1)	92*** (21.6)	370*** (149.1)
Nagpur	303*** (69.0)	209*** (57.8)	478*** (202.8)
Observations	28,056	10,320	17,736
Respondents	981	365	616

Notes: Conditional logits on the baseline earnings-arm choice sets pooled with the midline choice sets of control students. Destinations enter relative to the student's own home (the own district at baseline and home at midline); the salary attribute enters as level dummies relative to the expected level ($\pm 25\%$, $+50\%$), common across waves; wave-specific opt-out constants; standard errors clustered by student. Panel A reports the coefficient on a single indicator for any away option and the implied marginal utility per one percent of expected salary, averaged across the three realized level movements (delta-method standard errors). Panel B reports the marginal rate of substitution, the destination coefficient divided by the salary slope: the leave-home row pools all away options into one indicator, and the four destination rows come from the model with one indicator per destination. The men and women columns re-estimate each model on the gender subsample. Observations are alternatives (three per choice task, including the opt-out); respondents are students. The sample is unscreened; Appendix Table A10 repeats every cell on the attentive sample. Stars follow the p -value of the underlying utility coefficient. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 2: Effects on educational migration

	Index	Applied	Financing	Enrolled or plans to	Applied away	Migrated away	Migrated for education
Assigned to mentorship	0.155*** (0.054)	0.031** (0.014)	0.029 (0.027)	0.046** (0.021)	0.040 (0.026)	0.027 (0.022)	0.018 (0.021)
Control mean	0.000	0.906	0.521	0.788	0.335	0.216	0.181
Observations	1299	1391	1357	1383	1365	1403	1374

Notes: Each column reports the intention-to-treat coefficient from an OLS regression of the outcome on an indicator for assignment to the mentorship program, with robust standard errors in parentheses. The sample is all students with a completed endline interview; when no student interview is available, outcomes are taken from a proxy interview with a parent or another family member. A move away from home is any move beyond the student's home as defined in Section 2, including to the district capital; Appendix Table A25 reports the same outcomes for moves out of the home district. Each indicator counts a realized action or, for students who have not yet acted, a planned one within the following two months. *Applied*: applied or plans to apply for further studies. *Financing*: applied or plans to apply for a scholarship or loan to fund further studies. *Enrolled or plans to*: enrolled, or plans to enroll, in a degree college, vocational institute, or coaching program. *Applied away*: applied or plans to apply to at least one college away from home. *Migrated away*: has moved, or plans to move, away from home. *Migrated for education*: both migrated away and enrolled or planning to enroll. *Index*: inverse-covariance-weighted average of the six components (Anderson, 2008). Each component is standardized using the control-group mean and standard deviation; weights are the row sums of the inverted covariance matrix of the standardized components, estimated on control-group observations with all components observed. For respondents missing some components, the index aggregates the observed components only, with weights recomputed from the inverse of the covariance submatrix of the observed components and re-normalized to sum to one; the resulting index is re-standardized to mean zero and standard deviation one in the control group. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3: Applications and moves by destination

	All students			Outside a district capital		
	ITT	Control mean	N	ITT	Control mean	N
<i>Panel A. Applied or plans to apply to a college in ...</i>						
Home	-0.4 (2.7)	52.0	1391	-2.8 (3.1)	46.2	1056
District capital	4.9*** (1.8)	11.3	1391	6.9*** (2.4)	14.7	1056
Elsewhere in the home district	-0.5 (1.1)	4.8	1391	-1.2 (1.4)	5.7	1056
Pune or Mumbai	0.9 (1.5)	7.7	1391	1.0 (1.8)	9.0	1056
Elsewhere	-0.7 (1.4)	7.7	1391	-0.5 (1.8)	9.3	1056
<i>Panel B. Moved or plans to move for education to ...</i>						
District capital	2.1* (1.2)	4.0	1403	2.9* (1.5)	5.3	1068
Elsewhere in the home district	0.8 (1.3)	5.6	1403	1.6 (1.5)	5.6	1068
Pune or Mumbai	0.3 (1.3)	5.7	1403	-0.3 (1.6)	7.2	1068
Elsewhere	-0.8 (1.2)	5.9	1403	-0.3 (1.5)	6.9	1068

Notes: Each cell reports the intention-to-treat effect, in percentage points, on an indicator for having applied (Panel A) or moved (Panel B), realized or planned within the following two months, to a college in the destination named in the row; robust standard errors in parentheses; control means in percent. The sample and the definitions of applying and moving are those of Table 2; the right-hand columns restrict the sample to students whose home is outside a district capital (Ahilyanagar city or Sangli–Miraj), classified from the home reported at endline by the student or, failing that, by the proxy respondent or the school records. Panel A categories are not mutually exclusive, since a student can apply to colleges in several destinations. Panel B categories partition students who moved or plan to move away from home by destination. *District capital*: the capital of the student's home district. *Elsewhere in the home district*: any other taluka of the home district. *Pune or Mumbai*: the two metropolitan destinations in which the mentors study. *Elsewhere*: any other district. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 4: Effects on educational migration, by gender

	Men		Women		Difference (women – men)	N
	ITT	Control mean	ITT	Control mean		
Migration index (SD)	0.231** (0.098)	0.039	0.074 (0.072)	-0.024	-0.157 (0.121)	1115
<i>Components</i>						
Applied	0.024 (0.025)	0.913	0.027* (0.014)	0.951	0.003 (0.028)	1133
Financing	0.075 (0.046)	0.509	-0.018 (0.038)	0.624	-0.093 (0.060)	1136
Enrolled or plans to	-0.005 (0.035)	0.838	0.058** (0.030)	0.791	0.064 (0.046)	1123
Applied away	0.098** (0.047)	0.413	0.003 (0.036)	0.305	-0.095 (0.059)	1128
Migrated away	0.053 (0.043)	0.279	0.017 (0.029)	0.170	-0.036 (0.052)	1140
Migrated for education	0.003 (0.042)	0.259	0.031 (0.027)	0.132	0.029 (0.050)	1123
<i>By destination</i>						
Applied: district capital	0.091** (0.039)	0.178	0.024 (0.026)	0.117	-0.067 (0.047)	1133
Applied: Pune or Mumbai	0.042 (0.031)	0.100	-0.018 (0.022)	0.094	-0.060 (0.037)	1133
Migrated: district capital	0.042* (0.025)	0.059	0.012 (0.015)	0.031	-0.031 (0.029)	1140
Migrated: Pune or Mumbai	-0.002 (0.027)	0.090	-0.008 (0.016)	0.051	-0.006 (0.031)	1140

Notes: Direct student interviews only; proxy interviews are excluded because the proxy layer is not balanced on baseline characteristics within gender (Appendix I). For each outcome, the three estimates come from one OLS regression of the outcome on assignment to mentorship, an indicator for being female, and their interaction, with robust standard errors: the *Men* column reports the coefficient on assignment, the *Women* column the sum of that coefficient and the interaction (with the standard error of the sum), and the *Difference* column the interaction itself, whose stars test whether the effect differs by gender. The index is the inverse-covariance-weighted index of Table 2 re-estimated within the direct-interview sample and held fixed across genders, in control-group standard deviations; components are shares, defined as in Table 2; the destination rows are the application and move indicators of Table 3. Control means are for the control group of each gender. N: students in the regression. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 5: Effects on beliefs about destinations (endline)

	District capital		Pune	
	Treated $\times$ Away	N	Treated $\times$ Away	N
<i>Panel A. All students</i>				
Non-pecuniary	0.055 (0.119)	931	-0.135 (0.117)	975
Returns (salary, employment)	-0.006 (0.105)	931	0.021 (0.096)	975
Costs (monthly expenses)	-0.424*** (0.160)	891	0.006 (0.195)	929
<i>Panel B. Students outside a district capital</i>				
Non-pecuniary	0.057 (0.119)	929	-0.305* (0.175)	476
Returns (salary, employment)	-0.005 (0.105)	929	0.067 (0.147)	476
Costs (monthly expenses)	-0.428*** (0.160)	889	0.040 (0.287)	456
<i>Panel C. Outside a district capital, by gender</i>				
Non-pecuniary: men	0.123 (0.176)	407	-0.128 (0.277)	220
Non-pecuniary: women	0.006 (0.165)	522	-0.470** (0.220)	256
difference (women – men)	-0.117 (0.241)	929	-0.341 (0.353)	476
Returns (salary, employment): men	0.016 (0.188)	407	0.323 (0.214)	220
Returns (salary, employment): women	-0.010 (0.113)	522	-0.171 (0.200)	256
difference (women – men)	-0.026 (0.219)	929	-0.495* (0.293)	476
Costs (monthly expenses): men	-0.742*** (0.230)	391	0.205 (0.365)	210
Costs (monthly expenses): women	-0.132 (0.220)	498	-0.096 (0.436)	246
difference (women – men)	0.610* (0.318)	889	-0.301 (0.568)	456

Notes: Endline beliefs about the randomly assigned comparison destination, either the student's district capital or Pune, relative to beliefs about home. Each cell is the coefficient on Treated $\times$ Away from a regression, on the stacked home and comparison-destination observations of each student, of a belief index on assignment to mentorship, an indicator for the comparison destination, and their interaction, with gender $\times$ stream fixed effects (and comparison-city fixed effects in the district-capital columns), inverse-probability weights for response to the endline belief module, and standard errors clustered by student in parentheses. Indices are inverse-covariance-weighted averages (Anderson, 2008) standardized on the control group's home observations, so coefficients are in control-home standard deviations. *Non-pecuniary*: anticipated happiness, safety, mental wellbeing, whether people back home would think worse of the student, days per month with family, number of close friends, and parental approval, each signed so that higher is better. *Returns*: expected monthly salary after graduating and the probability of employment. *Costs*: expected monthly expenses (living costs plus one twelfth of yearly tuition, winsorized at the 99th percentile within treatment arm and destination), in their natural direction, so a positive coefficient means treated students expect the destination to cost more relative to home. Panels B and C restrict the sample to students whose home is outside a district capital. In Panel C the difference row reports the Treated $\times$ Away $\times$ Female coefficient from the fully interacted regression. N: student-destination observations. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.


A Appendix: Higher-education supply in Maharashtra

This appendix documents the concentration of higher-quality higher-education supply in Pune and the Mumbai region, where the mentors study, relative to the study districts (Ahilyanagar and Sangli), and provides the institutional context referenced in Section 2.


Supply across talukas. Figure A1 maps higher-education supply at the taluka level by entry-route group. Science and commerce-and-arts degree routes are concentrated in relatively few talukas, whether measured by the number of institutions (Panel A) or sanctioned intake (Panel B); non-degree vocational provision reaches more talukas but remains uneven on both measures.

Figure A1: Higher-education supply across Maharashtra talukas, by entry-route group.

A. Number of institutions



B. Reported sanctioned intake




Notes: The unit is a taluka. *Science* includes engineering, pharmacy, B.Sc. and GNM nursing, agriculture and allied undergraduate programs, and planning. *Commerce & Arts* includes management/BBA/BMS/BBM, BCA/integrated MCA, five-year law, design, fine art/AAC, hotel management, and integrated education; three-year law is excluded because it requires a prior degree. *Non-UG Vocational* includes DVET ITIs and active MSBSVET regular and STEP programs. Panel A counts distinct provisional campuses within each route group, linking records across sources where possible and using source-specific identifiers otherwise, so a campus can appear in more than one column if it offers routes in multiple groups; address-resolved records cover 1,119 of 1,380 Science campuses, 680 of 857 Commerce & Arts campuses, and 1,848 of 2,733 Non-UG Vocational campuses. Panel B sums reported sanctioned intake (nominal capacity, not current vacancies or seats available to a particular applicant), for which address-resolved campuses account for 90 %, 85 %, and 78 % of Science, Commerce & Arts, and Non-UG Vocational intake, respectively. Positive values are shown on a continuous logarithmic colour scale, with legend labels in original units; the three institution maps (Panel A) share one common scale and the three intake maps (Panel B) share a separate common scale. Gray denotes no mapped supply and can reflect either genuinely absent supply or an unresolved campus location. The two study districts, Ahilyanagar and Sangli, are outlined (solid and dashed, respectively). The latest available period is used for each route: 2025–26 for CET/CAP routes, 2026–27 for agriculture and DVET, and the active MSBSVET register on July 22, 2026. Sources: State Common Entrance Test Cell, Maharashtra; Maharashtra agricultural admissions records; Directorate of Vocational Education and Training; Maharashtra State Board of Skill, Vocational Education and Training; and Survey of India/Bharat Maps taluka boundaries archived with the project.

Centralized admissions and quotas. Engineering, pharmacy, and management seats in Maharashtra are allocated primarily through the Maharashtra Common Entrance Test (MHT-CET), a state-administered exam taken after twelfth grade. Allocation to specific programs is centralized: candidates submit ranked program preferences and are

assigned to seats based on their MHT-CET score, subject to category- and state-level quotas that reserve a substantial share of seats for in-state applicants. Medical and elite-engineering seats are allocated through separate national exams (NEET and JEE, respectively). The figures below use MHT-CET 2024 administrative data; the patterns are similar for pharmacy and management programs.


Seats per eligible candidate. Among engineering programs in the top quartile of admission cutoffs, Pune and the Mumbai region offer markedly more seats per eligible MHT-CET candidate than Ahilyanagar and Sangli (Figure A2). A home-district candidate who clears the top-quartile cutoff therefore faces a thinner local option set than they would in Pune or Mumbai.

Figure A2: Engineering seats per eligible candidate, MHT-CET 2024.



Distribution of admission cutoffs. Origin districts also skew toward programs with lower admission cutoffs (Figure A3). The right tail of the cutoff distribution (the most selective programs) is concentrated in Pune and the Mumbai region.

Figure A3: Distribution of minimum MHT-CET scores by district. Origin districts have fewer, less-competitive programs.



Observational income gaps. Table A1 reports the income, consumption, and education gaps between Pune and the Mumbai region and the rest of Maharashtra in the India Human Development Survey, the source of the household-survey figures quoted in Section 2.

Table A1: Observational income and education gaps within Maharashtra (IHDS-II).

	Income PC		Consumption PC		Education	
	All HH (1)	HH with graduate (2)	All HH (3)	HH with graduate (4)	Years (5)	Graduate or higher (6)
Pune and the Mumbai region	2.01*** (0.12)	1.60*** (0.27)	1.02*** (0.06)	0.85*** (0.13)	3.06*** (0.14)	0.10*** (0.01)
Mean other Maharashtra	1.91	3.13	1.80	2.49	6.80	0.07
N	2,877	592	2,873	592	8,631	8,631

Notes: Each cell reports the coefficient on an indicator for Pune and the Mumbai region from a separate weighted regression on the IHDS-II 2011–12 Maharashtra cross-section; the Mumbai region here comprises Mumbai and Thane (Mumbai Suburban is not a separate IHDS district). Income and consumption are monthly per-capita values in thousands of 2012 INR, adjusted for spatial price differences using an HCES 2023–24 Fisher unit-value index for Pune, the Mumbai region, and the rest of the state, by rural and urban sector. *Graduate or higher* is at least 15 completed years of education. Household regressions (columns 1–4) report robust standard errors; individual education regressions (columns 5–6) are restricted to ages 18–64, include age fixed effects, and cluster standard errors by household. Sampling weights are applied but are not calibrated for district-level representativeness, so the estimates are descriptive conditional differences, not representative district estimates or causal returns to migration (sorting, household composition, labor-market selection, and survey coverage all contribute). *Mean other Maharashtra* is the weighted outcome mean in the estimation sample outside these two regions. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

B Appendix: Mentor applicant pool and migrant-peer evidence

This appendix describes the pool of students who applied to be Campus Trail mentors. These data are not meant to be representative of educational migrants, nor of the mentors who ultimately worked in the program. Applicants selected into the form because they were interested in mentoring, and the experiences survey was completed by a smaller

subset of this pool. I therefore use the evidence here descriptively: as anecdotal information on what kinds of destination-specific experience the program could transmit to mentees, rather than as a comparison group for the student sample.

Applicant profile. Table A2 summarizes the applicant pool characteristics. Applicants are older and academically selected relative to the high-school students in the experiment. Most study in Pune or Mumbai and most report having moved specifically for studies, usually alone. They report substantially higher living costs compared to the student sample, given the selection into migration and application to become a mentor, I do not interpret these as a representative estimate of what the marginal student in the experiment would spend if they moved.

Table A2: Summary statistics for the mentor applicant pool

Variable	N	Mean	SD	Median
<i>Application form applicant pool</i>				
Female	83	0.59	0.49	1.00
Current undergraduate	83	0.53	0.50	1.00
Studying in Pune/Mumbai region	83	0.88	0.33	1.00
Moved specifically for studies	83	0.78	0.41	1.00
Moved alone	83	0.78	0.41	1.00
Knew someone in destination	83	0.53	0.50	1.00
12th-board percentage	78	79.26	12.68	80.50
<i>Experiences survey respondent pool</i>				
Monthly expenses (Rs.)	33	14,948	10,257	13,000
Annual tuition (Rs.)	33	75,047	79,090	40,000
Accommodation plus food (Rs./month)	33	16,369	18,435	10,000
No financial assistance	35	0.60	0.50	1.00

Notes: The application-form panel uses applicant records with a non-missing institution name. The experiences-survey panel uses respondents with a non-missing institution name. Indicator means are shares within this selected pool. Monetary variables are cleaned by dropping negative values and values of Rs. 1,000,000 or more, which are treated as invalid/sentinel entries. The table is descriptive only; the applicant pool is not representative of educational migrants or of the mentors who ultimately worked in the program.

What destination contacts provide. The experiences survey asks migrant applicants whether they had contacts in the destination before moving, what help those contacts provided, and what help would have been useful if they had no contacts. Across both groups, the recurring themes are practical and destination-specific: what living in the city costs, where a student can live safely and affordably, which colleges or programs

have a good reputation, what campus life is like, and how to manage the first weeks after arrival.

What applicants propose to provide to mentees. The application form also asks applicants which topics they would discuss in a one-hour conversation with a twelfth-standard student considering study in their city. In addition, each applicant receives a randomized vignette describing one concern a hypothetical mentee has: costs and housing, academic returns and jobs, missing home and making friends, family and community pressure, daily adjustment, or safety and navigation, and is asked how they would advise the student.

Analyzing the responses offer rich insights into how mentor pool views the transition into college life, and what type of information, support or encouragement could flow to mentees. When the vignette is about costs and housing, applicants often translate the concern into a concrete budget: rent ranges, food costs, transport passes, hostels, PGs, tiffin services, scholarships, educational loans, and sometimes part-time work. When the vignette is about academic returns, they tend to compare the destination option with staying closer to home, discussing course structure, faculty, placements, alumni, internships, and the skills students would need beyond the degree itself. When the concern is family or community pressure, applicants often switch to a parent-facing argument: moving is framed as an investment in career prospects, exposure, independence, and long-run family benefit rather than as a break from home. For homesickness and daily adjustment, the advice is more relational. Applicants normalize loneliness, describe routines for managing cooking, budgeting, and chores, and point to roommates, seniors, clubs, student groups, professors, and counselling services as supports. For safety and navigation, they usually combine reassurance with concrete precautions: safe areas, transport options, emergency contacts, campus security, and people the student could call after arrival.

C Appendix: Destination options by survey wave

Several survey modules ask students to report beliefs, preferences, information, or hypothetical choices for named destinations. Across waves, the common design goal is to measure how students view other destinations relative to home. The destination set changes across waves because the definition of home is refined from district to taluka, because the intervention introduces a mentor city, and because the endline phone survey requires a shorter comparison set. The table refers to these destination-specific modules;

questions about actual or planned applications, enrollment, work, and migration could record broader locations.

Table A3: Destination options used in destination-specific survey modules

Survey wave	Destination pool	Student-specific destinations evaluated	Wave-specific notes
Baseline	Own district; Pune; Mumbai; Nagpur.	Each student evaluated three destinations: own district plus two of Pune, Mumbai, and Nagpur. The first comparison destination was randomized between Pune and Mumbai. The second was randomized between Nagpur and the other of Pune/Mumbai.	Pune, Mumbai, and Nagpur were selected as the three prominent higher-education destinations in Maharashtra. Home-relative outcomes compare each away destination to the student's own district.
Midline	Own taluka; district capital; Pune; Mumbai.	Each student evaluated own taluka plus two comparison destinations. The set always included the student's mentor city: the actual mentor city for treated students and a pseudo-mentor city for comparison students, assigned using the same matching procedure. For students not from the district capital, the remaining comparison destination was randomized between the district capital and the other of Pune and Mumbai. For students from the district capital, the two comparison destinations were Pune and Mumbai.	District capital was introduced after baseline because district-level "home" questions could be interpreted as the district capital rather than the student's own taluka. Home-relative outcomes compare away destinations to own taluka.
Endline	Own taluka; Pune; district capital.	The fixed comparison destination was assigned before endline. Students from the district-capital taluka were assigned Pune. Other students were randomized within treatment, gender, and stream cells to either Pune or the district capital, with roughly one third assigned to Pune and two thirds to the district capital. The phone survey first elicited the realized or planned destination, which could be outside this fixed comparison set, and then asked the additional home/comparison modules needed to form the pre-specified belief differences.	The shorter phone-survey design preserves the same home-relative logic for two focal comparisons: Pune versus home and district capital versus home.

Notes: "District capital" refers to Ahilyanagar city for students from Ahilyanagar district and Sangli city/Sangli-Miraj for students from Sangli district. At baseline, home is measured at the district level; at midline and endline, home is measured at the taluka level.

D Appendix: Discrete choice experiment

This appendix documents the construction and presentation of the DCE attributes described in Section 4.

Overview. In both survey waves, each choice task presented two hypothetical college profiles and a "No university" opt-out. The randomized profiles varied in destination and in one return or quality attribute, while the remaining profile details were fixed at the student level after being constructed from the student's own survey responses. The

designs were generated in Stata using the `dcreate` module for efficient discrete-choice designs (Hole, 2015).

Destination sets. The DCE used the student-specific destination set constructed for the relevant survey wave. Appendix C describes the baseline and midline destination pools and selection rules.

Profile construction. Several profile values were anchored to earlier survey responses, while others were fixed at the student level to keep the tradeoff focused on destination and the arm-specific return or quality measure. Table A4 summarizes both the anchor source and the displayed construction for each profile item.

Table A4: DCE profile attributes and anchors

Wave/arm	Profile item	Fixed or randomized	Anchor/source	Displayed construction
Baseline and midline	Field of study	Fixed at student level	“What would you like to study at university?”; if multiple fields were selected, the survey asked for the most preferred field.	All profiles shown to the student were labelled as colleges in the student’s most-preferred field.
Baseline and midline	Destination	Randomized	Student-specific destination set constructed for the wave; see Appendix C.	Destination was randomized within the student’s set.
Baseline and midline, earnings arm	Expected earnings	Randomized	“How much money do you think you will earn per year at age 25?” The survey elicited a bracket and an exact amount; the DCE used the bracket anchor.	0.75, 1.00, 1.25, or 1.50 times the bracket anchor, converted to rupees.
Baseline, employment arm	Employment probability	Randomized	0–10 response to “How likely is it that you will be working at a full-time job when you are 25 years old?” The response was bounded at 2–8 and multiplied by 10.	Anchor –10, anchor, anchor +10, or anchor +20 percentage points.
Baseline, cutoff arm	Admission cutoff	Randomized	Expected entrance-exam score when relevant; otherwise expected 12th-board score. The anchor was bounded between 20 and 80. Program-level MHT-CET data were used to convert expected scores into cutoff scores corresponding roughly to program-rank movements.	Students saw cutoff scores out of 100. In practice, these scores corresponded roughly to 50 ranks better, the reference cutoff, or 50 ranks worse; see implementation note below.
Baseline and midline	Monthly room and food costs	Fixed at student level	Family budget for monthly living expenses while studying in college, including room, food, and other expenses.	90 percent of the lower endpoint of the selected budget bracket; Rs. 2,000 if the bracket minimum was zero. Fixed across all profiles shown to the student.
Baseline and midline	Yearly college fees	Fixed at student level	Family budget for yearly college fees.	70 percent of the lower endpoint of the selected budget bracket; Rs. 3,000 if the bracket minimum was zero. Fixed across all profiles shown to the student.

Notes: A lower numerical program rank corresponds to a better-ranked and more selective program. Monetary amounts were displayed in rupees.

Baseline design. The baseline DCE has three randomized arms: expected earnings, employment probability, and admission cutoff. In each arm, the intended candidate set was the full factorial of three destinations and four levels of the arm-specific attribute shown in Table A4.

For each arm, I generated 70 D-efficient two-profile choice sets using uninformative priors and blocked them into ten seven-question blocks. Students were assigned to one of the resulting 30 blocks. Motivated by Clayton et al. (2023), I then added an eighth task to each block by repeating the second task with the left and right profiles mirrored. This mirrored pair provides a robustness check for consistency and side bias. Within each block, both question order and profile order were randomized. The opt-out was displayed in the baseline survey as a third response option alongside the two randomized college profiles.

Cutoff-arm implementation note. Due to a coding error in the baseline survey, the cutoff score corresponding to the intended 100-rank-better level was replaced by the reference cutoff score. The realized baseline cutoff arm therefore contained displayed cutoff scores corresponding roughly to three rank positions: 50 ranks better, the reference cutoff, and 50 ranks worse.

Midline design. The midline DCE focuses on the earnings arm and uses priors estimated from the baseline earnings-arm choices. Specifically, I estimate conditional-logit priors using baseline respondents who gave consistent answers on the mirrored baseline pair, and then use those priors to generate a Bayesian D-efficient design. The midline design includes three destinations and the same four earnings levels as the baseline earnings arm. The design contains 80 choice sets blocked into ten eight-question blocks, with question order and profile order randomized within block.

Pooled estimation sample. The estimates in Table 1 pool the baseline earnings-arm choice sets with the midline choice sets of control students, so that the district capital, offered only at midline, can be placed on the same ladder as the out-of-district cities. Every destination is coded relative to the student's own home: the own district at baseline and home at midline. For baseline students who live in a district capital the own-district option is their home, and their Pune, Mumbai, and Nagpur options identify those destination coefficients; baseline students who live elsewhere in the district were never offered their district capital as a separate option, so their away options enter with destination-specific nuisance indicators and contribute only to the salary coefficient. At

midline, students living outside a capital identify the district-capital coefficient, and all control students contribute to the Pune and Mumbai coefficients. The salary levels are common across waves, each wave and residence group has its own opt-out constant, and standard errors are clustered by student. The main-text estimates use every completed choice task (981 students, 555 at baseline and 614 at midline, 9352 choice sets); Table A10 repeats them after dropping, in both waves, students who chose the same option position in every task. The belief-controlled estimates add the eight destination-specific beliefs of Section 4, measured in the wave of the choice task (the baseline city module for baseline tasks, the midline city module for midline tasks), with moving difficulty set to its definitional home value in both waves; each wave's beliefs enter with their own coefficients. The no-university alternative carries the reference level of the randomized attributes and, in the belief-controlled models, the student's beliefs about home, since not enrolling means staying there; Appendix Table A8 reports the alternatives.

Presentation. Before the baseline choice tasks, students saw the note: "Please click on the "Play Button" to watch the video for instructions for the next section!" The accompanying instruction image began: "Which college would you choose? You're about to see information about pairs of imaginary colleges. In each pair, imagine you've scored high enough to get admission into both." It then explained that, for each college, the survey would show living costs, college fees, location, and the arm-specific quality measure: "placement packages" in the earnings arm, "share of students who got jobs through campus placement" in the employment arm, and "cutoffs" in the cutoff arm. The screen listed the four profile details and then stated: "The only differences are the 4 details we'll show you. Everything else about the colleges is the same!" It ended: "In each round, you'll be given only three options: College 1, College 2, or not going to college at all. Which one would you choose? There's no right or wrong answer – just what feels perfect for your future!" Figure A4a shows the instruction poster and an example choice task.




Figure A4: DCE instruction and choice-task screens shown in the baseline survey.

D.1 Additional results

Table A5 reports the baseline away penalty (at baseline every away option lies outside the home district) separately by numeraire arm, on the attentive sample, together with the implied willingness to pay in each arm's units. Table A6 prices the pooled destination ladder of Table 1 in percentage points of the probability of employment, using the baseline employment arm in place of the earnings arm, and Table A7 prices it in rupees of yearly salary rather than in percent of each student's own expectation. Figure A5 shows the salary-level coefficients with and without belief controls, Table A8 the belief-controlled leave-home penalty by sample and by opt-out convention, and Table A9 the item-level decomposition of the pooled leave-home penalty. Table A10 repeats every main-text cell on the attentive sample.

Table A5: The migration penalty: baseline DCE estimates

	Earnings	Employment	Cutoffs
Away from home	-0.817 (0.068)	-0.938 (0.070)	-0.804 (0.073)
Attribute slope (per unit)	0.0065 (0.0013)	0.0302 (0.0032)	0.0035 (0.0007)
No university	-1.932 (0.114)	-1.853 (0.107)	-1.893 (0.104)
MRS to stay home	126 (27)	31 (4)	231 (53)
men	74 (16)	26 (5)	249 (125)
women	233 (113)	35 (6)	223 (56)
MRS units	% of earnings	pp employment	cutoff percentiles
Nr. students	489	510	475

Notes: Conditional logits on choice sets of three destination options, estimated on the attentive sample (dropping students who chose the same option position in every task), standard errors clustered by student. Each arm column uses that arm's randomized quality attribute, entered as level dummies relative to the expected level: $\pm 25\%$ and $+50\%$ for earnings; -10 , $+10$, $+20$ pp for placement; ± 50 ranks for cutoffs, shown as the corresponding admission-cutoff scores.^a The attribute slope is the implied marginal utility per unit of the numeraire (per percent of earnings, percentage point of placement, or program rank for cutoffs), averaged across the realized level movements assuming linearity; delta-method standard errors. The MRS to stay home is the ratio of the away coefficient to this slope. The men and women rows re-estimate the arm's model on each gender subsample. The cutoffs MRS is in program ranks; as the text explains, students appear to treat a higher cutoff as admission difficulty rather than quality, so this column is reported for completeness rather than read as a quality tradeoff.

^a A display error in the survey software made the fourth cutoff level (100 points better than the reference) show the reference value; those alternatives are coded at the reference level and the cutoff slope averages the two realized movements.

Table A6: The migration penalty in employment terms

	All	Men	Women
<i>MRS: percentage-point increase in $P(\text{employment})$ to migrate</i>			
Away from home	24.2*** (3.4)	19.8*** (4.9)	26.7*** (4.7)
District capital	7.7** (3.2)	7.5* (4.6)	7.9* (4.3)
Pune	17.4*** (3.0)	16.8*** (4.6)	17.9*** (4.0)
Mumbai	37.0*** (5.0)	28.0*** (6.9)	41.5*** (6.9)
Nagpur	43.7*** (9.3)	60.5** (31.2)	42.6*** (10.2)
Students	959	358	601


Notes: Construction as in Table 1, with the baseline employment arm in place of the earnings arm: the same midline control choice sets supply the district-capital coefficient, their salary levels entering as nuisance attributes, and the MRS divides each destination coefficient by the placement slope, averaged over the -10 , $+10$, and $+20$ percentage-point movements. Unscreened sample. Stars follow the p -value of the underlying utility coefficient. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A7: The migration penalty in rupees

	All	Men	Women
<i>Panel A. Utility coefficients</i>			
Away from home	-0.597*** (0.055)	-0.543*** (0.086)	-0.625*** (0.070)
Salary (per lakh Rs per year)	0.084 (0.011)	0.082 (0.016)	0.086 (0.016)
<i>Panel B. MRS: lakh Rs per year of salary to migrate</i>			
Away from home	7.13*** (1.12)	6.59*** (1.54)	7.48*** (1.64)
District capital	2.16** (1.01)	2.42* (1.50)	1.99 (1.34)
Pune	5.11*** (0.97)	5.76*** (1.51)	4.76*** (1.27)
Mumbai	10.72*** (1.64)	8.55*** (2.11)	11.78*** (2.43)
Nagpur	16.45*** (3.28)	19.17*** (5.69)	15.20*** (4.08)
Observations	28,056	10,320	17,736
Respondents	981	365	616

Notes: Sample and specification as in Table 1, with the salary attribute entered in rupees rather than as level dummies: each displayed salary is the level's multiplier (0.75, 1, 1.25, 1.5) times the student's own expected yearly earnings, so the attribute is the rupee change relative to the student's expectation, in lakh (100,000) rupees per year, and enters with a single coefficient. Panel A reports the coefficient on a single indicator for any away option and the coefficient on the rupee attribute. Panel B reports the MRS, the destination coefficient divided by the rupee coefficient, in lakh rupees of yearly salary; because the movement is in a common currency, the MRS is comparable across students with different expected earnings, whereas the percentage MRS of Table 1 scales with each student's own anchor. The leave-home row pools all away options; the destination rows come from the model with one indicator per destination; the men and women columns re-estimate each model on the gender subsample. Observations are alternatives (three per choice task, including the opt-out); respondents are students. Stars follow the p -value of the underlying utility coefficient. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Figure A5: The numeraire with and without belief controls: salary-level coefficients



Notes: Coefficients on the salary levels relative to the expected salary in the pooled model of panel (c) of Figure 3, without controls (gray) and with the eight destination-specific beliefs of each wave (blue); their average slope per one percent of expected salary is the numeraire that converts destination coefficients into the MRS. Whiskers are 95 % confidence intervals, standard errors clustered by student. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A8: Belief controls and the leave-home penalty across samples

	Away coefficient (clogit)			MRS, % of salary			OLS decomposition, % absorbed				
	No beliefs	Beliefs	% absorbed	No beliefs	Beliefs	% change	Total	Emotional	Social	Practical	Students
Panel A. Samples (opt-out row carries home beliefs)											
All arms and midline control (paper)	-0.709	-0.170	76.0	152	37	75.8	61.3	13.3	14.1	34.0	1,682
Baseline, all arms	-0.789	-0.229	71.0	144	41	71.6	57.3	14.7	14.3	28.4	1,682
Baseline, earnings arm	-0.774	-0.448	42.2	142	79	44.5	35.9	15.8	12.6	7.5	555
Baseline, employment arm	-0.847	-0.203	76.0	—	—	—	60.0	13.2	13.8	33.0	575
Baseline, cutoffs arm	-0.740	-0.038	94.8	—	—	—	75.1	14.0	18.1	42.9	552
Midline control only	-0.484	-0.022	95.4	129	6	95.4	76.2	7.2	12.6	56.3	614
Earnings arm and midline control, student-relative (Table 1)	-0.587	-0.196	66.7	129	42	67.3	54.3	10.3	12.3	31.6	981
Panel B. Opt-out row convention (all arms and midline control)											
Home beliefs (paper)	-0.709	-0.170	76.0	152	37	75.8	61.3	13.3	14.1	34.0	1,682
Mean of the two colleges' beliefs	-0.709	-0.410	42.1	152	86	43.3	46.2	14.7	16.9	14.6	1,682

Notes: Each row re-estimates the leave-home model, a single indicator for any away option, without and with the eight destination-specific beliefs of each wave, on the choice sets with complete belief data. Panel A varies the sample. The first six rows use the pooled construction of panel (c) of Figure 3: own district is home for every baseline student, destination coefficients are common across arms and waves, and each arm has its own attribute block and opt-out constant. The last row of Panel A is the student-relative construction of Table 1. Panel B holds the pooled sample fixed and varies what the no-university alternative carries in the belief block. In the paper it carries the student's beliefs about home, since not enrolling means staying there; the alternative is the average of the two colleges' beliefs, which makes the opt-out row neutral so that the belief coefficients are identified only from the choice between places. Columns 1-3: the conditional-logit coefficient without and with beliefs and its percentage reduction. Columns 4-6: the MRS, the coefficient divided by the salary slope, in percent of expected salary, and its percentage reduction; the slope is identified only where the earnings arm is present, so the employment and cutoff rows have no MRS. Columns 7-10: the linear-probability decomposition with choice-set fixed effects, in which the total absorbed equals the sum of the three family shares exactly (Gelbach, 2016). Standard errors are omitted; the pooled row's bootstrap intervals are in Table A9.

Table A9: Decomposition of the pooled leave-home penalty by belief

	All	Men	Women
<i>Emotional well-being</i>	13.3	16.2	12.5
95% CI	[10.0, 16.7]	[9.6, 23.1]	[8.9, 16.3]
Happiness	12.6	16.0	11.3
Mental wellbeing	0.7	0.2	1.2
<i>Social ties</i>	14.1	11.7	15.1
95% CI	[10.7, 18.0]	[6.4, 17.9]	[10.1, 20.5]
Reputation back home	0.4	0.2	0.5
Days seeing family	0.9	1.0	0.7
Close friends at destination	2.3	2.0	2.5
Parental approval	10.5	8.6	11.4
<i>Practical and lived experience</i>	34.0	13.7	46.5
95% CI	[24.1, 44.6]	[-2.0, 30.3]	[33.5, 59.7]
Safety	4.2	-1.6	7.4
Moving difficulty	29.8	15.3	39.1
Total absorbed (%)	61.3	41.6	74.1
95% CI	[51.6, 72.4]	[27.8, 58.6]	[61.1, 88.0]
Away coefficient, no beliefs (OLS)	-0.279 (0.015)	-0.246 (0.023)	-0.300 (0.020)
Away coefficient, with beliefs (OLS)	-0.108 (0.021)	-0.144 (0.031)	-0.078 (0.028)
Students	1682	659	1023

Notes: Linear-probability version of the leave-home model on the pooled sample of panel (c) of Figure 3 (all three baseline arms and the midline control choice sets, own district as home at baseline), with choice-set fixed effects, estimated on the complete-belief choice sets, without and with the eight destination-specific beliefs of each wave. Each item's share is the belief's coefficient in the model with beliefs multiplied by the within-choice-set gap in that belief between the away option and home, summed over the two waves' versions of the item, and expressed as a share (%) of the away coefficient without beliefs (Gelbach, 2016); because the auxiliary regressions carry the same controls as the short model, the item shares sum exactly to the total absorbed. Family subtotals and the total carry 95% confidence intervals from a cluster bootstrap that resamples students, stratified by gender (1000 replications). Men and women columns re-estimate the decomposition on each gender subsample. Because the beliefs are positively correlated, the joint model divides shared covariation among them, so single-item shares understate their standalone importance.

Table A10: The migration penalty on the attentive sample

	Main text	Attentive sample
<i>MRS (% of expected salary)</i>		
Away from home	129*** (26.3)	123*** (24.0)
Out of district	142*** (29.3)	137*** (27.1)
District capital	39** (19.1)	34** (17.7)
Pune	93*** (21.9)	88*** (20.2)
Mumbai	196*** (39.1)	192*** (36.5)
Nagpur	303*** (69.0)	259*** (58.3)
Away from home, men	72*** (16.0)	65*** (14.0)
Away from home, women	227*** (92.8)	227*** (90.7)
<i>Beliefs and the penalty (belief sample)</i>		
Away MRS, belief-controlled	42** (18.4)	51*** (16.9)
Change in MRS (%)	67.3	58.3
Beliefs absorbed, OLS (%)	61.3	54.1
Students	981	862
Students, belief sample	981	862

Notes: The main-text column repeats the estimates of Table 1, the belief-controlled leave-home MRS on the same earnings-arm build, and the pooled-sample decomposition of Table A9. The attentive-sample column re-estimates every cell after dropping, in both survey waves, students who chose the same option position in every choice task, the pre-specified attention check. Change in MRS is the percentage reduction in the leave-home MRS of the earnings-arm build when the belief controls are added; beliefs absorbed is the reduction in the linear-probability away coefficient on the pooled sample of all baseline arms and midline control. Standard errors clustered by student; stars follow the p -value of the underlying utility coefficient. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

D.2 Inattention and measurement error

Random inattention in conjoint experiments attenuates estimated effects toward zero (Clayton et al., 2023). Following their design recommendation, the second and eighth choice tasks mirror each other: the same profiles reappear with positions swapped. A respondent answering attentively should choose the same profile both times.

Between 72 and 76 % of students do (Table A11). Under the benchmark model in which a respondent answers randomly with probability 2ϵ (so two independent measurements of the same choice disagree with probability $2\epsilon(1 - \epsilon)$), the implied swapping error is $\epsilon = 0.14\text{--}0.17$, and observed effects understate true effects by a factor of $1 - 2\epsilon \approx 0.67\text{--}0.71$. The table reports three checks. First, the away effect on the choice probability, rescaled by $1/(1 - 2\epsilon)$, rises by roughly half in magnitude; the correction scales all attribute effects equally, so the MRS is unaffected by this form of error. Second, re-estimating on the consistent respondents only leaves the away coefficient and the implied MRS close to the full-sample estimates. Third, dropping the repeated task changes the coefficients by less than one standard error. The migration penalty is therefore not an artifact of inattentive responding, and its conversion into MRS units is robust to it by construction.

Table A11: Inattention in the DCE: mirrored-task consistency and corrections

	Earnings	Employment	Cutoffs
Consistent on mirrored task (IRR)	0.777	0.788	0.754
Implied swapping error ϵ	0.128	0.120	0.144
Implied attenuation $1 - 2\epsilon$	0.744	0.759	0.712
Away effect on choice (AMCE)	-0.136 (0.026)	-0.176 (0.026)	-0.022 (0.024)
AMCE corrected by $1/(1 - 2\epsilon)$	-0.183 (0.034)	-0.232 (0.034)	-0.031 (0.034)
Away coefficient, consistent only	-0.918 (0.084)	-1.110 (0.088)	-0.901 (0.092)
MRS to stay home, consistent only	136 (34)	32 (4)	206 (46)
Away coefficient, excl. repeated task	-0.822 (0.069)	-0.945 (0.072)	-0.806 (0.074)
Nr. students (consistent / all)	380 / 489	402 / 510	358 / 475

Notes: IRR = share of students choosing the same profile in the mirrored second and eighth tasks. The swapping error solves $2\epsilon(1 - \epsilon) = 1 - \text{IRR}$, treating the three-option choice as binary between the two college profiles. The AMCE is the coefficient on the away indicator from a linear probability model of choice on the attributes with choice-set fixed effects; the corrected row rescales it by $1/(1 - 2\epsilon)$ per Clayton et al. (2023). Consistent-only rows re-estimate the conditional logit of Table A5 on students who passed the mirror check; the final coefficient row drops the repeated (eighth) task. Standard errors clustered by student.

E Appendix: Belief elicitation questions

At baseline, students answer the questions below about each of the three destinations they are shown: their own district and two others (Appendix C). In the wording, [des-


destination] is the city being asked about and [home] is the student's own district. In the analysis every item is recoded so that higher values are better (safer, happier, more approving, easier to move, and so on).

Table A12: Baseline non-monetary belief questions

Belief	Question	Response scale
Happiness	If you live and study in [destination], how happy do you think you would be?	0–10 (0 = completely unhappy, 10 = completely happy)
Mental well-being	If you live and study in [destination], how likely is it that you might face difficulties with your mental wellbeing?	0–10 (0 = definitely would not, 10 = definitely would)
Safety	How safe would you feel living and studying in [destination]?	0–10 (0 = completely unsafe, 10 = completely safe)
Reputation back home	How likely is it that other people in [home] will talk badly about you, if you were living and studying in [destination]?	0–10 (0 = would never, 10 = would definitely)
Parental approval	All things considered, would your parents approve of you living and studying in [destination]?	0–10 (0 = would not approve at all, 10 = would completely approve)
Days seeing family	If you live and study in [destination], how many days per month would you meet family and friends in [home]?	Count, 0 to 40+ days
Close friends	If you live and study in [destination], how many close friends do you think you would have in the first year of studying there?	Count, 0 to more than 10
Moving difficulty	Imagine you are moving to [destination] for college (finding accommodation, registering, locating healthcare and food, learning to navigate the city and transport). How challenging would it be to move to [destination] for college?	Four points: not challenging at all, slightly, very, overwhelmingly challenging

Notes: Each item is asked separately for each of the three elicited destinations.

Figure A6: Beliefs about non-pecuniary outcomes by destination, by gender



Notes: Same construction as Figure 4, with the wave-adjusted destination means computed separately for men (teal) and women (violet): the student's own town, the capital of the student's own district, Pune, Mumbai, and Nagpur, pooling the baseline city module (all students) and the midline control group. Items are in their own units; mental well-being is the reverse-coded probability of mental-health difficulties, reputation the probability that people back home would not think worse of the student, and moving ease reverses the four-point difficulty item with a definitional home value. Freedom, asked only at baseline, is not shown. The gender columns of Table A13 report the corresponding away-home gaps in standard-deviation units.

Table A13: Baseline beliefs about destinations: home versus away

			Away – home (home SDs)		
	Home mean	Away mean	All	Men	Women
<i>Emotional well-being</i>					
Happiness	7.37	5.75	-0.53*** (0.028)	-0.54*** (0.043)	-0.53*** (0.037)
Mental wellbeing	0.63	0.55	-0.22*** (0.025)	-0.15*** (0.039)	-0.26*** (0.032)
<i>Social ties</i>					
Reputation back home	0.66	0.61	-0.14*** (0.022)	-0.04 (0.034)	-0.20*** (0.029)
Days seeing family	9.04	7.32	-0.18*** (0.019)	-0.19*** (0.031)	-0.18*** (0.023)
Close friends at destination	3.67	2.67	-0.28*** (0.019)	-0.27*** (0.032)	-0.29*** (0.024)
Parental approval	6.98	4.96	-0.62*** (0.026)	-0.46*** (0.041)	-0.72*** (0.033)
<i>Practical and lived experience</i>					
Safety	7.43	5.28	-0.72*** (0.028)	-0.60*** (0.046)	-0.79*** (0.036)
Moving difficulty ^a	3.00	1.59	-1.67*** (0.020)	-1.50*** (0.033)	-1.79*** (0.026)


Notes: Home and away means in the item's own units (0–10 scales, probabilities, or counts). The gap columns are coefficients from regressions of the item (standardized by its home mean and home standard deviation) on an away indicator, pooled across destinations (at baseline every away destination lies outside the home district), overall and by gender; standard errors clustered by student. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

^a Moving difficulty is asked only about away destinations; the home value is set to “not challenging at all,” its definitional value, and the item is standardized by the away-sample standard deviation (the home value has no variance).

Information items

Figure A7 reports the distribution of the baseline items on navigating college admissions and on familiarity with the destination cities that underlie Fact 3 of Section 4.

Figure A7: Baseline knowledge about college admissions and about destination cities.



Notes: Each bar shows the distribution of responses to one baseline item, oriented from least (*Not at all*) to most (*A great deal*). Panel A: “How difficult is it for you to [find and compare college information; understand admission requirements; make progress on college decisions; find someone to ask for advice]?”, where *Not at all* denotes “very difficult, I don’t know how to do it” and *A great deal* “not difficult at all, I know how to do it.” Panel B, pooled over the destination cities each student was asked about: “How much do you know about studying and student life in [city]?” (*Knows the destination city*) and the perceived ease of the move (*Moving there feels manageable*, the reverse of “how challenging would moving there be”). Students report low absolute levels on every item.

F Appendix: College rating exercise

This appendix documents the college rating exercise used to measure how students value college location when profile attributes are varied directly. It was administered at baseline and re-administered at midline; the design below describes the midline version. The baseline version was identical in structure except that students rated three profiles rather than five, and the destination was the student’s own district, displayed as the district-capital city, or one of the same distant cities used in the discrete choice experiment (Pune, Mumbai, or Nagpur). The willingness-to-pay benchmark reported in Section 4 uses the baseline wave.

Overview. Students rated five hypothetical college profiles one at a time. Each profile was described as a college in the student’s most-preferred field. Each screen asked how likely the student would be to accept admission if the displayed college were the only option. Responses were recorded on a 0–10 scale.

Profile construction. Each profile was generated inside the survey form using fresh random draws. The destination was drawn from the student’s midline destination list, described in Appendix C. Fees, living costs, hostel information, admission cutoff, placement share, average earnings, and higher-education share were randomized at the profile level. Earnings and higher-education shares were not drawn independently of the rest of the profile: profiles with higher placement shares tended to have higher earnings and lower higher-education shares, with additional random variation. These correlations made the profiles look more like realistic college information while preserving experimental variation in each attribute.

Table A14: College rating profile attributes

Profile item	Fixed or randomized	Construction
Field of study	Fixed	Profiles were labelled as colleges in the student’s most-preferred field.
Destination	Randomized	One destination was drawn from the student-specific midline destination list. In the analysis, displayed destinations are coded into the location indicators used in the paper, including away from home, out of district, and mentor city.
Yearly college fees	Randomized	Rs. 1,000 plus a uniform draw between Rs. 0 and Rs. 100,000, rounded to the nearest rupee.
Monthly food and living costs	Randomized	Rs. 5,000 plus a uniform draw between Rs. 0 and Rs. 50,000, rounded to the nearest rupee.
Hostel information	Randomized	Four equally likely categories: clean and safe college hostel; college hostel available but quality uncertain; only private hostel or PG options available; and no information.
Admission cutoff	Randomized	A uniform draw between 40 and 100, displayed as the cutoff out of 100 for the general category.
Job placement share	Randomized	A uniform draw between 0 and 100, displayed as the share of students placed in jobs after graduation.
Average yearly earnings	Randomized	Generated as a positive function of the placement share plus random noise, with a lower bound of Rs. 100,000 per year.
Higher-education share	Randomized	Generated as a negative function of the placement share plus random noise, bounded between 0 and 100.


Notes: Monetary values were displayed in rupees. All randomized values were generated separately for each profile.

Presentation. Before the rating screens, students saw the note: “Please click on the “Play Button” to watch the video for instructions for the next section!” The same screen displayed an instruction poster and an instruction video. Each profile then appeared as a small table with the destination, yearly fees, monthly food and living costs, hostel information, admission cutoff, share selected for higher education, share placed in jobs, and average yearly earnings.

The rating prompt read: “If this were your only college option, how likely would you be to accept admission? Rate from 0-10, where:” The left endpoint defined 0 as “I would definitely NOT attend this college (0% chance),” and the right endpoint defined 10 as “I

would definitely attend this college (100% chance).” The survey hint below the response scale read: “Rate from 0-10.” Figure A8a shows the instruction poster and an example profile rating screen.

(a) Instruction poster



Commerce College	
City	Mumbai
Fees (yearly)	₹65,479
Food and living expenses (monthly)	₹46,342
Hostel	College hostel available, quality uncertain
Cutoff for admission (for General Category)	65 out of 100
After graduation..	
28% (out of 100%) of students were selected for further education	
18% (out of 100%) of students were placed in jobs	
Average earnings (yearly) ₹639,988	

If this were your only college option, how likely would you be to accept admission? Rate from 0-10, where:

0 means "I would definitely **NOT** attend this college (0% chance)" 10 means "I would definitely attend this college (100% chance)"

Rate from 0-10.

0 1 2 3 4 5 6 7 8 9 10

(b) Profile rating screen

Figure A8: College rating exercise screens shown in the midline survey.

F.1 Willingness to pay to stay home

I recover a willingness to pay to remain at home from the baseline rating exercise with a hedonic regression of the 0–10 profile rating on an away indicator and the randomized profile attributes, using the same destination definition as the baseline discrete choice experiment: the home profile is the student’s own district, displayed as the district-capital city, and the away profiles are Pune, Mumbai, or Nagpur. Because ratings respond to the fee attribute but only weakly to the earnings attribute, I read the away penalty off the tuition slope, as the annual fee reduction that would leave a student indifferent, and then express it as a share of the student’s own expected earnings so that it can be compared with the discrete-choice willingness to pay in Section 4. Table A15 reports the estimates overall and by gender.

Table A15: Migration penalty and willingness to pay in the baseline college-rating exercise

	All	Men	Women
Away penalty (rating points)	-1.025*** (0.119)	-0.804*** (0.188)	-1.166*** (0.154)
Yearly tuition (per Rs 1,000)	-0.0119*** (0.0019)	-0.0105*** (0.0030)	-0.0125*** (0.0024)
Monthly living cost (per Rs 1,000)	-0.0120*** (0.0036)	-0.0159*** (0.0059)	-0.0097** (0.0046)
Placement share (per pp)	0.0063** (0.0028)	0.0035 (0.0045)	0.0083** (0.0035)
Yearly earnings (per Rs 1,000)	0.0003** (0.0001)	0.0005*** (0.0002)	0.0001 (0.0002)
Admission cutoff (per point)	-0.0052* (0.0029)	-0.0046 (0.0048)	-0.0060 (0.0037)
Higher-education share (per pp)	0.0076 (0.0091)	0.0298** (0.0140)	-0.0063 (0.0118)
Clean and safe hostel	0.427*** (0.153)	0.395 (0.246)	0.433** (0.194)
Hostel, quality uncertain	0.118 (0.150)	0.130 (0.239)	0.080 (0.194)
Only private / PG hostel	0.076 (0.153)	0.058 (0.239)	0.086 (0.197)
Implied WTP to stay home (Rs '000/yr)	86.0*** (17.2)	76.6*** (27.8)	93.0*** (22.7)
as % of median expected earnings	86	51	133
Memo: DCE penalty (% of earnings)	126	74	233
Median expected earnings (Rs '000/yr)	100.0	150.0	70.0
Nr. students	1474	575	899
Nr. profiles	4422	1725	2697

Notes: Hedonic OLS of the 0–10 college-profile rating on an away indicator and the randomized profile attributes (yearly tuition, monthly living costs, placement share, yearly earnings, admission cutoff, higher-education share, and hostel type), estimated on the attentive sample (students who did not straightline the discrete choice experiment, matching the DCE): 1,474 students, three profiles each. The home profile is the student's own district, displayed as the district-capital city, and away is Pune, Mumbai, or Nagpur, as in the baseline discrete choice experiment. The away penalty is the rating gap; the continuous attributes enter as marginal rating effects per Rs 1,000 (tuition, living cost, earnings), per percentage point (placement share, higher-education share), or per cutoff point (admission cutoff), and the three hostel indicators are relative to a profile giving no hostel information. Ratings respond to fees far more than to the earnings attribute, so the implied willingness to pay to stay home is recovered in tuition units: the away penalty divided by the tuition slope (the annual fee reduction that offsets it), reported in '000 INR and as a percentage of the group's median expected earnings. The admission-cutoff slope is negative, so students rate more-selective programs lower even though the cutoff here was drawn independently of their own expected score, corroborating the pattern discussed in Section 4; and consistent with the safety concerns documented there, a clean and safe hostel is the amenity students value most, raising the rating well above every other amenity level. The memo row reproduces the DCE migration penalty (% of expected earnings) from Table A5. The men and women columns re-estimate the regression on each subsample. Standard errors clustered by student in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

G Appendix: Sample selection and corrections

This appendix details the restrictions applied to the pool of shortlisted applicants to arrive at the randomization sample (Section 5), and collects the descriptive tables on the schools, the randomization balance, and selection into take-up.

G.1 Sample restrictions

I exclude schools with fewer than five interested applicants for Campus Trail. The sampling and data-collection strategy is designed to recruit students genuinely interested in the program; as an additional data-quality measure, I conduct pre-specified interest and attention checks and keep only students who pass them. Specifically, I exclude students who:

- respond “No” to whether they are interested in applying to Campus Trail both at the beginning and at the end of the survey;
- respond “I want to stop studying after Class 12” when asked about educational plans, or “I am not interested in college recommendations” or “I am not interested in meeting a mentor” when asked about potential destinations and topics for mentor guidance; or
- are flagged by field enumerators as inattentive (a score below 3 on the 0–10 attentiveness scale), or explicitly marked as such.

G.2 Schools, balance, and take-up

Appendix Figure A9 compares the 2026 HSC board results of the study-sample schools with district and state benchmarks; Table A16 compares the study schools with the other secondary schools in the two study districts; Table A17 reports baseline balance across treatment arms; and Table A18 reports which baseline characteristics predict take-up of mentorship among treated students.

Figure A9: HSC Board Exam results (2026): study-sample schools relative to district and state benchmarks. Study-sample schools are slightly lagging.




Table A16: School comparison: study sample vs. other schools in the study districts

	Study sample Mean/(SD)	Other schools Mean/(SD)	Difference (district FE)
Higher-secondary board: State	0.980 (0.140)	0.954 (0.211)	0.017 (0.021)
Rural school	0.804 (0.401)	0.683 (0.466)	0.140** (0.061)
Marathi medium	0.902 (0.300)	0.765 (0.424)	0.121*** (0.045)
Has internet	0.941 (0.238)	0.943 (0.232)	-0.007 (0.035)
Has electricity	1.000 (0.000)	0.996 (0.067)	0.005 (0.003)
Private school	0.196 (0.401)	0.364 (0.481)	-0.138** (0.059)
No. of 12th-grade students	127.216 (94.006)	132.157 (178.396)	-7.710 (15.221)
Share girls in 12th grade	0.486 (0.147)	0.411 (0.229)	0.065*** (0.023)
Teachers per student	0.029 (0.009)	0.056 (0.141)	-0.025*** (0.005)
Classrooms per student	0.025 (0.010)	0.091 (0.461)	-0.056*** (0.015)
Projectors per student	0.005 (0.005)	0.013 (0.062)	-0.009*** (0.003)
Joint <i>F</i> -test <i>p</i>		0.000	
Schools	51	668	719

Notes: Study-sample schools versus all other secondary schools in the two study districts (Ahilyanagar and Sangli), from UDISE administrative school records. One study-sample school is unmatched in these records, so the comparison covers 51 of the 52 study schools. Cells report group means with standard deviations below; the Difference column reports the study-sample minus other-schools difference from a regression with district fixed effects, with robust standard errors below, and the final row reports the *p*-value of a joint *F*-test. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A17: Summary statistics and randomization balance

	Control Mean/(SD)	Treated Mean/(SD)	Difference (T–C)
Female	0.617 (0.486)	0.599 (0.490)	-0.018 (0.024)
Age	17.221 (0.617)	17.259 (0.637)	0.038 (0.031)
General category	0.367 (0.482)	0.355 (0.479)	-0.012 (0.023)
Past board marks	71.005 (11.334)	70.795 (10.926)	-0.210 (0.543)
Household size	5.424 (2.268)	5.383 (2.206)	-0.041 (0.109)
Asset index	-0.000 (1.000)	0.028 (1.151)	0.028 (0.053)
Has electricity	0.971 (0.168)	0.956 (0.204)	-0.014 (0.009)
Tap water at home	0.508 (0.500)	0.535 (0.499)	0.027 (0.024)
Pukka house	0.647 (0.478)	0.684 (0.465)	0.037 (0.023)
Mother completed class 10	0.537 (0.499)	0.519 (0.500)	-0.018 (0.024)
Father completed class 10	0.687 (0.464)	0.697 (0.460)	0.010 (0.023)
Travels once a year or more	0.748 (0.435)	0.723 (0.448)	-0.025 (0.022)
Internet access	0.826 (0.379)	0.817 (0.387)	-0.009 (0.019)
Self-efficacy	16.218 (5.032)	16.424 (4.727)	0.205 (0.238)
Joint orthogonality F -test p		0.400	
Observations		1,682	

Notes: Baseline means and standard deviations by treatment arm for the analysis sample ($N = 1,682$). The final column reports the treatment–control difference; the final row reports the p -value of an F -test of joint orthogonality. Standard errors are robust.

Table A18: Selection into take-up

	No take-up Mean/(SD)	Took up Mean/(SD)	Difference (up–no)
Female	0.619 (0.486)	0.576 (0.495)	-0.043 (0.034)
Science stream	0.650 (0.477)	0.734 (0.442)	0.084*** (0.032)
Age	17.334 (0.631)	17.171 (0.633)	-0.163*** (0.044)
General category	0.341 (0.475)	0.371 (0.484)	0.030 (0.033)
Past board marks	69.753 (10.914)	72.018 (10.827)	2.264*** (0.759)
Expected board marks	72.227 (9.628)	73.303 (9.846)	1.076 (0.680)
Asset index	0.076 (1.152)	-0.028 (1.148)	-0.104 (0.080)
Father completed class 10	0.697 (0.460)	0.697 (0.460)	0.000 (0.032)
Internet access	0.805 (0.397)	0.832 (0.375)	0.027 (0.027)
Household size	5.516 (2.315)	5.226 (2.063)	-0.289* (0.152)
Willingness to take risk	5.430 (3.141)	5.797 (3.083)	0.367* (0.217)
Number of close friends	4.545 (3.501)	4.424 (3.343)	-0.121 (0.239)
Self-efficacy	15.982 (4.679)	16.942 (4.737)	0.960*** (0.329)
Has decided on a plan	1.020 (0.910)	1.082 (0.890)	0.061 (0.063)
Any network away from home	0.558 (0.497)	0.576 (0.495)	0.018 (0.035)
P(apply) away from home	0.413 (0.284)	0.487 (0.287)	0.075*** (0.020)
Rating of away destinations	5.656 (2.952)	6.042 (2.865)	0.386* (0.203)
Social desirability (M–C)	6.827 (2.312)	6.858 (2.350)	0.031 (0.163)
Effort index	-0.019 (1.003)	0.033 (0.985)	0.052 (0.069)
Knowledge index	-0.047 (0.903)	0.021 (1.055)	0.068 (0.069)
Joint F -test p		0.000	
Treated students	446	380	826

Notes: Baseline characteristics of treated students who took up mentorship (at least fifteen minutes of recorded conversation) versus those who did not. Columns report group means with standard deviations below; the Difference column reports the took-up minus did-not-take-up difference from an OLS regression, with the robust standard error below, and the final row reports the p -value of a joint F -test that the characteristics jointly predict take-up. The covariate set is the one used to predict take-up in the causal-forest analysis.
* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

H Appendix: Parent survey and parent–student comparisons

Data collection. The parent instrument mirrors the student baseline (Section 5.1): educational aspirations and expectations for the child, affordability of tuition and monthly expenses, a discrete choice experiment over hypothetical college profiles (four choice sets, same three numeraire arms as the student DCE), ratings of the three destinations in the student’s elicitation set, a two-city belief module covering home and the nearer of the student’s two elicited cities (city order randomized), networks, destination consideration, and the gender-attitude and family-decision batteries. The parent sample over-represents more-connected households: baseline covariates jointly predict which parents were reached ($p = 0.003$), though this selection does not differ by treatment arm ($p = 0.94$; Table A19).

Table A19: Selection into the parent survey

	Predicts response (selective)	Differential by arm (Treated $\times$ X)
Female	0.048* (0.025)	-0.069 (0.050)
Age	-0.004 (0.012)	-0.003 (0.024)
General category	0.005 (0.025)	-0.007 (0.051)
Science stream	-0.002 (0.026)	-0.010 (0.052)
Past board marks	0.011 (0.012)	-0.021 (0.024)
Asset index	-0.031** (0.012)	0.035 (0.024)
Household size	-0.015 (0.012)	-0.016 (0.024)
Mother completed class 10	-0.038 (0.024)	-0.041 (0.049)
Father completed class 10	-0.005 (0.026)	0.002 (0.053)
Internet access	-0.047 (0.032)	-0.028 (0.064)
Self-efficacy	0.030** (0.012)	-0.006 (0.024)
Any network away from home	-0.020 (0.025)	-0.006 (0.049)
P(apply) away from home	0.031** (0.012)	-0.010 (0.024)
Joint test p	0.003	0.938
Observations		1,682

Notes: Selection into the parent survey on the full experimental sample ($N = 1,682$). The *Predicts response* column reports, for each baseline characteristic, the coefficient from a regression of an indicator for the student's parent being surveyed on that characteristic (selective attrition); the *Differential by arm* column reports the coefficient on the characteristic's interaction with treatment assignment (whether that selection differs by arm). Continuous predictors are standardized, so their coefficients are per standard deviation; binary predictors are in levels. Baseline characteristics jointly predict which parents were reached ($p = 0.003$), but this selection does not differ by treatment arm ($p = 0.94$). Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Parents' migration penalty is a multiple of their children's. Estimating the same conditional logit as in Section 4 on the parent DCE, the average parent would forgo roughly a 450 % increase in the child's post-graduation earnings to keep them at home, versus

roughly 150 % for students when the student experiment is estimated with the same specification on the full baseline sample (the pooled-wave estimate of Section 4 is 129%; employment arm: 168 pp versus 31 pp). In the cutoff arm the comparison is made in coefficient units: the parent away coefficient is three times the student one (-2.28 vs. -0.74), but parents barely respond to the cutoff attribute, so no meaningful rank-denominated MRS exists for them; parent-arm estimates are in any case imprecise, with ≈ 270 parents and four choice sets per arm. The destination ratings tell the same story in simpler form: the rating penalty for away destinations is -5.4 points (0–10) for parents against -1.5 for students, driven both by parents rating home higher (9.2 vs. 7.6) and away cities lower (3.8 vs. 6.1).

Parents perceive a larger wedge on every non-monetary belief. Parents expect away destinations to compare worse to home than students do on each elicited item: expected happiness falls by 4.1 points (0–10) for parents versus 1.2 for students, safety by 4.7 versus 2.0, days seeing family by 9.6 versus 1.5 per month, and parental approval by 4.7 versus 1.9.

Students misperceive their parents' main objection. Both sides are asked the *main* reason the parent would object to the child attending college in the distant city, from the same reason list. Only 13.8 % of matched pairs name the same reason. The disagreement is systematic: parents overwhelmingly name safety (34.6 % of parents, versus 7.8 % of students guessing safety), while students over-attribute concerns about housing (14.3 % vs. 5.9 %) and jobs (10.2 % vs. 0.9 %). Both sides agree on expense (26.1 % vs. 22.5 %).

Parental attitudes and program engagement. Parents' willingness for the child to migrate (listing any away destination when asked where the child should look for colleges: 44 %; approval of the child living in the nearer elicited city: 4.3 on a 0–10 scale) predicts mentorship take-up only weakly on its own and adds nothing once the student's own baseline application intentions are controlled for (parent measures joint $p = 0.75$): engagement with the mentor is driven by the student's own willingness.

I Appendix: Endline data collection and attrition

This appendix details the endline survey implementation, the tracking protocol used to limit attrition, and the associated weighting and bounding procedures summarized in Section 5.1.

I.1 Survey implementation

Endline surveys are conducted over the phone and take on average twenty minutes; respondents are offered 150 INR (about 1.6 USD) for completing the interview. Whenever a student is unavailable, the field team conducts a proxy interview with a parent, a friend, or a headmaster or teacher at the student’s school, collecting information on the primary outcomes only. Proxy interviews take on average five minutes and are not incentivized. The primary analysis relies on data collected directly from students; when only a proxy interview is available, proxy-reported outcomes enter the main analysis with an indicator for proxy-sourced observations, and I report robustness checks that exclude them. Where the same outcome is available from both the student and a proxy, the proxy data are used only for robustness.

I.2 Tracking and weighting

Endline data collection begins with a standard phone-tracking phase in which enumerators attempt to reach every student using all available numbers (the student’s own and those of parents or other household members) with repeated attempts at different times of day and attempts to obtain updated contact information from parents and schools.

After this phase I construct the intensive-tracking frame. A student is classified as “not reached” if, by a pre-specified tracking checkpoint, no student or proxy interview has been completed. The intensive-tracking unit is the school, the natural cluster for tracking students: students were recruited through their schools, and schools hold updated contact information through records, teachers, and classmates. Letting N_s denote the number of not-reached students in school s at the checkpoint, schools with at least five not-reached students are eligible for intensive tracking.¹⁵ I select up to ten schools using sequential probability-proportional-to-size sampling without replacement, drawing each school with probability proportional to its number of not-reached students among schools not yet selected. In selected schools the field team visits in person and attempts to trace all not-reached students, completing in-person student interviews where possible and proxy interviews otherwise.

For each school I compute the inclusion probability π_s implied by this procedure by simulation, conditional on the realized N_s . I report unweighted estimates using all completed surveys and weighted estimates that account for the unequal probability of

¹⁵I exclude the smallest schools (those with fewer than five not-reached students) up to at most 5% of the attrited sample.

intensive tracking, with weights

$$w_i = \begin{cases} 1, & \text{if } i \text{ was surveyed before the tracking checkpoint,} \\ 1/\pi_s, & \text{if } i \text{ was surveyed through intensive tracking.} \end{cases}$$

I report overall and arm-specific attrition at the checkpoint and after intensive tracking, test for differential attrition by treatment arm, and test whether baseline covariates predict attrition and do so differentially by treatment. Table A21 reports endline response by treatment arm on the full experimental sample, and Table A22 reports, for both survey waves, whether baseline characteristics predict response (selective attrition) and whether that selection differs by arm (differential attrition). The midline was administered in person during school visits and involved no tracking: enumerators surveyed the students who were present, so midline non-response is selective on who attended school that day but does not differ by treatment arm. The endline, reached by phone with the protocol above, shows mild differential attrition concentrated among treated students with the strongest baseline migration intentions. Table A20 reports response by arm and wave under the sample definitions used in the paper.

Table A20: Survey response and attrition, by treatment arm

	Control	Treated	Difference
<i>Panel A. Midline (in school, untracked)</i>			
Surveyed	0.717	0.714	-0.003 (0.022)
<i>Panel B. Endline (by phone, tracked)</i>			
Any interview	0.848	0.838	-0.010 (0.018)
Direct student interview	0.674	0.685	0.011 (0.023)
Proxy interview only	0.174	0.153	-0.022 (0.018)
Full primary data (student)	0.686	0.708	0.022 (0.022)
Full primary data (any source)	0.791	0.797	0.006 (0.020)
Observations		1,682	

Notes: Share of the full experimental sample responding at each wave, by treatment arm, with the treatment–control difference from an OLS regression and its robust standard error below. Panel A reports the midline response rate; the midline was administered in schools without tracking, so response reflects who attended on the survey day. Panel B reports endline response under alternative sample definitions: *Any interview* (a completed student or proxy interview), *Direct student interview*, *Proxy interview only*, and *Full primary data* observed from the student or from any source. Endline fieldwork is ongoing and these rates are not final until the data freeze. Selective attrition (whether baseline characteristics predict response) and differential attrition (whether that selection differs by arm) are reported in Table A22. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A21: Endline response and estimates under intensive-tracking weights

Panel A. Survey response, by treatment arm			
	Control	Treated	Difference
Any interview at the checkpoint	0.757	0.735	-0.022 (0.021)
Any interview	0.848	0.838	-0.010 (0.018)
Student interview	0.674	0.685	0.011 (0.023)
Proxy interview	0.174	0.153	-0.022 (0.018)
Full primary (student)	0.686	0.708	0.022 (0.022)
Full primary (any)	0.791	0.797	0.006 (0.020)
Recovered through intensive tracking Observations	92 of 109 pursued 1682		
Panel B. Primary outcomes, with and without intensive-tracking weights			
	Control mean	Unweighted	Weighted
Out-of-district index	-0.000	0.099* (0.053)	0.078 (0.056)
Migrated out of district	0.120	-0.000 (0.017)	0.007 (0.017)
Away-from-home index	0.000	0.155*** (0.054)	0.125** (0.056)
Migrated away from home	0.216	0.027 (0.022)	0.035 (0.022)

Notes: Panel A is the full experimental sample. Each row reports, for each arm, the share of students meeting the row definition; the Difference column reports the OLS coefficient on treatment assignment, with the robust standard error in parentheses below. *Any interview at the checkpoint*: a completed student or proxy interview by the tracking checkpoint, the point at which the not-reached frame was fixed and the intensive-tracking schools drawn. *Any interview*: a completed student or proxy interview as of the current extract, and so inclusive of students recovered since the checkpoint. *Student interview*: a completed interview with the student directly. *Proxy interview*: a completed proxy interview (parent, friend, or teacher) with no completed student interview. *Full primary (student)*: all six primary-outcome components observed from the student's own report, including partially completed interviews. *Full primary (any)*: all six components observed from any source. The recovery line reports how many of the not-reached students pursued in the drawn schools have been interviewed so far.

Panel B reports the primary outcomes on students with a completed interview, first unweighted and then weighted by the pre-specified intensive-tracking weight defined above: one for students surveyed before the tracking checkpoint, and $1/\pi_s$ for students recovered through intensive tracking, where π_s is the simulated probability that their school was drawn. A recovered student therefore stands in for the comparable not-reached students in the eligible schools that were not drawn. Because intensive tracking is in progress, the weighted column rests on the students recovered so far and moves further from the unweighted column as recovery proceeds; the two indices are inverse-covariance-weighted (Anderson, 2008). Not-reached students in schools ineligible for intensive tracking have no probability of recovery and are represented by neither column; they are assigned worst-case values in the bounds reported below. Response rates and estimates update as fieldwork proceeds and are not final until the data freeze. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A22: Selective and differential attrition at midline and endline

	Midline (in school, untracked)		Endline (phone, tracked)	
	Selective (X)	Differential ($T \times X$)	Selective (X)	Differential ($T \times X$)
<i>Baseline demographics</i>				
Female	0.104*** (0.023)	0.008 (0.046)	0.029 (0.018)	-0.014 (0.037)
Past board marks (SD)	0.029** (0.011)	0.016 (0.023)	0.016* (0.009)	0.000 (0.018)
Self-efficacy (SD)	0.018 (0.011)	-0.023 (0.022)	0.028*** (0.009)	-0.037* (0.019)
Asset index (SD)	-0.016 (0.012)	0.039* (0.023)	0.003 (0.008)	0.008 (0.017)
Household size (SD)	0.023** (0.011)	-0.029 (0.022)	0.005 (0.009)	-0.019 (0.019)
Internet access	-0.042 (0.028)	-0.033 (0.056)	-0.004 (0.023)	-0.064 (0.046)
Mother class 10	0.004 (0.022)	-0.012 (0.044)	0.025 (0.018)	0.033 (0.036)
Father class 10	-0.020 (0.024)	-0.066 (0.047)	0.024 (0.020)	0.036 (0.039)
Science stream	0.015 (0.024)	0.043 (0.048)	0.016 (0.019)	0.071* (0.039)
General social category	-0.009 (0.023)	0.087* (0.046)	0.021 (0.018)	0.057 (0.036)
Sangli district	-0.030 (0.022)	0.048 (0.044)	-0.053*** (0.018)	0.029 (0.036)
Lives in a district capital	-0.042 (0.026)	-0.006 (0.052)	0.005 (0.020)	0.009 (0.041)
<i>Baseline analogues of primary outcomes</i>				
Wants to live elsewhere at 25	0.002 (0.022)	-0.017 (0.044)	0.035* (0.018)	0.091** (0.036)
P(apply) away from home (SD)	0.009 (0.011)	0.021 (0.022)	0.013 (0.009)	0.008 (0.018)
Rating of away destinations (SD)	0.010 (0.011)	0.014 (0.022)	0.008 (0.009)	-0.031* (0.018)
Any network away from home	-0.005 (0.029)	-0.038 (0.058)	0.019 (0.023)	0.045 (0.047)
Joint test p	0.001	0.362	0.107	0.010
Observations		1682		

Notes: Each wave's two columns report OLS regressions of a response indicator on baseline characteristics, estimated on the full 1,682-student experimental sample. The *Selective* column regresses the indicator for being surveyed on the baseline characteristic alone; its coefficient shows whether the characteristic predicts response. The *Differential* column regresses response on the characteristic, treatment assignment, and their interaction, and reports the interaction coefficient; it shows whether selection into response differs by arm. Continuous predictors are standardized to control-group units, so their coefficients are per standard deviation; binary predictors are in levels. Midline response is a completed midline interview; endline response is a completed student or proxy interview. The midline was administered in person during school visits without tracking, so midline response reflects who attended school on the survey day; the endline was conducted by phone under the tracking protocol described above. "Joint test p " is the p -value of an F -test that all baseline characteristics jointly predict response (*Selective* columns) or that all treatment interactions are jointly zero (*Differential* columns). Endline response rates are not final until the data freeze. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

I.3 Attrition bounds

I report three sets of bounds following Coppock et al. (2017). First, for the binary primary outcomes I report double-sampling worst-case bounds, using the randomized intensive-tracking sample to estimate outcomes among initially not-reached students and assigning worst-case values to cases that remain unrecovered. Second, I report sensitivity bounds in which the never-surveyed are assumed to differ from the recovered by a range of magnitudes calibrated to the observed gap between students reached during standard phone tracking and students reached through intensive tracking. Third, I report post-stratified double-sampling bounds, defining poststrata by baseline gender and stream and aggregating poststratum-specific treatment-control differences using the baseline share of each poststratum, collapsing to gender only if the gender-by-stream cells are too sparse. Students in schools ineligible for intensive tracking have zero recovery probability and are assigned worst-case values. I also report Lee bounds (Lee, 2009) as a complementary approach.

Table A23 reports these procedures for the migration index. Intensive-tracking weights lower the point estimate for any move away from home from 0.15 to 0.12 standard deviations. The double-sampling worst-case and poststratified bounds are uninformative, because every student who remains unrecovered, 16% of the sample, is assigned a worst-case value; excluding students who explicitly refused the interview, the worst-case interval is 0.08 to 0.19.

Table A23: Attrition adjustments and bounds for the migration index

	Out of district	Away from home
Point estimate (complete case)	0.099	0.155
Intensive-tracking weighted	0.078	0.125
Double-sampling worst case	[-0.482, 0.620]	[-0.382, 0.611]
Sensitivity, missing shifted by the tracking gap:		
calibration gap g	0.008	-0.053
both arms at $+g$ (point)	0.078	0.125
both arms at $+2g$ (point)	0.078	0.126
adversarial, $\pm g$	[0.076, 0.080]	[0.111, 0.138]
adversarial, $\pm 2g$	[0.074, 0.082]	[0.098, 0.151]
Poststratified (gender $\times$ stream)	[-0.484, 0.612]	[-0.386, 0.600]
Poststratified (gender only)	[-0.484, 0.619]	[-0.384, 0.609]
<i>Deviations from the PAP (2026-09-01)</i>		
Worst case, refusals excluded (190 students)	[0.020, 0.136]	[0.081, 0.186]
Unrecovered at control prediction	0.069	0.109
(bootstrap s.e.)	(0.051)	(0.050)

Notes: Intention-to-treat estimates for the pre-specified migration index under the attrition procedures of the pre-analysis plan, for moves out of the home district (first column) and any move away from home (second column, the definition of Table 2). *Point estimate*: complete-case estimate on all students with a completed interview. *Intensive-tracking weighted*: weights of one for students surveyed before the tracking checkpoint and $1/\pi_s$ for students recovered through intensive tracking. *Double-sampling worst case*: bounds that assign the worst-case value of the index to students who remain unrecovered. *Sensitivity*: the never-surveyed are assumed to differ from the recovered by the calibration gap g , the observed difference between students reached in standard tracking and students reached through intensive tracking, in the same direction in both arms (point) or adversarially across arms (intervals), at one and two times g . *Poststratified*: double-sampling bounds aggregated within gender $\times$ stream or gender poststrata using baseline shares. The final block departs from the plan: worst-case bounds that exclude students who explicitly refused the interview, and an estimate that assigns each unrecovered student the control-group prediction at their own baseline characteristics, with a bootstrap standard error over the imputation and estimation. Estimates update as fieldwork proceeds and are not final until the data freeze.

I.4 Robustness to the primary-index definition

Table A24 re-estimates the migration index under alternative constructions and samples: the pre-specified index on all students with a completed interview; the same index with post-double-selection lasso controls; direct student interviews only; an expanded sample that adds resolved cases missing some components together with incomplete surveys that reached the primary module; and an index built from realized outcomes only, dropping planned actions from every component. Every index is re-standardized within its own estimation sample, so control means are zero by construction and the columns are directly comparable in control standard deviations. The effect is present in every column, between 0.13 and 0.20 standard deviations; excluding proxy interviews lowers it slightly and counting only realized outcomes raises it.

Table A24: Migration index under alternative definitions and samples

	Pre-specified	PDS controls	Direct interviews	Incl. incomplete	Realized outcomes only
Assigned to mentorship	0.155*** (0.054)	0.159*** (0.053)	0.141** (0.058)	0.133*** (0.050)	0.197*** (0.065)
Control mean	0.000	0.000	0.000	-0.000	0.000
Observations	1299	1299	1115	1485	1110

Notes: Each cell is the intention-to-treat coefficient on assignment to mentorship for the inverse-covariance-weighted migration index (Anderson, 2008), any move away from home, under one construction, with the standard error below. *Pre-specified*: the six-component index on all students with a completed interview, as in Table 2. *PDS controls*: adds post-double-selection lasso controls chosen from the baseline covariate set, with analytic standard errors. *Direct interviews*: restricts to direct student interviews, as in Appendix Table A26. *Incl. incomplete*: the expanded index, adding resolved cases missing some components together with incomplete surveys that reached the primary module. *Realized outcomes only*: drops planned actions from every component and counts only applications made, financing applied for, enrollment started, and moves completed by the interview date, on direct student interviews. Each index is re-standardized within its own estimation sample, so control means are zero by construction. Robust standard errors except in the PDS column. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

I.5 Spillovers between students

Because randomization was at the individual level within schools, treated students could in principle relay information to control classmates, which would attenuate the intention-to-treat estimates. I test for such interference using the best-friend nominations collected at midline, which identify, for most students, a specific classmate in the experimental sample. Control students whose best friend was assigned to treatment are no more likely to migrate or to apply to colleges away from home: the coefficient on having a treated best friend is -0.04 (s.e. 0.05) for any move away from home, and the estimates for the other primary outcomes are similarly small and, if anything, negative. The result is unchanged when I restrict to friendships of at least a year's standing, which predate the intervention and so cannot reflect any post-treatment change in whom a student names as a best friend (-0.04 , s.e. 0.06). I find no evidence that spillovers to control students attenuate the estimates.

J Appendix: Additional treatment-effect estimates

This appendix reports the estimates referred to in Section 7: the out-of-district version of the main migration table, the main table on direct student interviews only, and the split by whether the student's home is in a district capital.

Table A25: Effects on out-of-district educational migration

	Index	Applied	Financing	Enrolled or plans to	Applied away	Migrated away	Migrated for education
<i>Panel A. All students (student or proxy interview)</i>							
Assigned to mentorship	0.099* (0.053)	0.031** (0.014)	0.029 (0.027)	0.046** (0.021)	0.002 (0.021)	-0.000 (0.017)	-0.002 (0.016)
Control mean	-0.000	0.906	0.521	0.788	0.173	0.120	0.098
Observations	1299	1391	1357	1383	1365	1403	1374
<i>Panel B. Direct student interviews</i>							
Assigned to mentorship	0.077 (0.058)	0.024* (0.013)	0.017 (0.029)	0.033 (0.023)	0.003 (0.023)	-0.001 (0.019)	0.001 (0.018)
Control mean	0.000	0.937	0.579	0.809	0.182	0.122	0.099
Observations	1115	1133	1136	1123	1128	1140	1123

Notes: As Table 2, with a move defined as leaving the home district rather than leaving home: *Applied away* counts applications to colleges outside the home district, *Migrated away* moves out of the home district, and *Migrated for education* both. *Applied*, *Financing*, and *Enrolled or plans to* are unchanged. The index is the inverse-covariance-weighted average of the six out-of-district components, constructed as in Table 2. Panel A uses all students with a completed endline interview (student or proxy); Panel B restricts to direct student interviews, with the index re-estimated within that sample. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A26: Effects on educational migration, direct student interviews only

	Index	Applied	Financing	Enrolled or plans to	Applied away	Migrated away	Migrated for education
Assigned to mentorship	0.141** (0.058)	0.024* (0.013)	0.017 (0.029)	0.033 (0.023)	0.047 (0.029)	0.036 (0.025)	0.024 (0.024)
Control mean	0.000	0.937	0.579	0.809	0.346	0.212	0.181
Observations	1115	1133	1136	1123	1128	1140	1123

Notes: As Table 2, restricted to students with a completed direct interview; proxy-sourced observations are excluded and the index is the inverse-covariance-weighted average re-estimated within the direct-interview sample. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A27: Effects on educational migration, by district-capital residency

	Index	Applied	Financing	Enrolled or plans to	Applied away	Migrated away	Migrated for education
<i>Panel A. Students living outside a district capital</i>							
Assigned to mentorship	0.210*** (0.061)	0.031* (0.017)	0.071** (0.031)	0.044* (0.025)	0.059* (0.031)	0.045 (0.027)	0.030 (0.026)
Control mean	-0.000	0.905	0.527	0.773	0.406	0.252	0.212
Observations	985	1056	1031	1050	1038	1068	1043
<i>Panel B. Students living in a district capital</i>							
Assigned to mentorship	-0.002 (0.110)	0.032 (0.029)	-0.098* (0.055)	0.051 (0.038)	0.003 (0.033)	-0.015 (0.031)	-0.007 (0.029)
Control mean	0.000	0.909	0.500	0.836	0.099	0.097	0.079
Observations	314	335	326	333	327	335	331

Notes: As Table 2, estimated separately for students whose home is outside a district capital (Panel A) and students living in one (Ahilyanagar city or Sangli-Miraj; Panel B), classified from the home reported at endline by the student or, failing that, by the proxy respondent or the school records. Within each panel the index is the inverse-covariance-weighted average of the six components re-estimated on that panel's sample (complete cases, standardized on its control group), so control means are zero by construction and the two panels are on their own scales. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.